



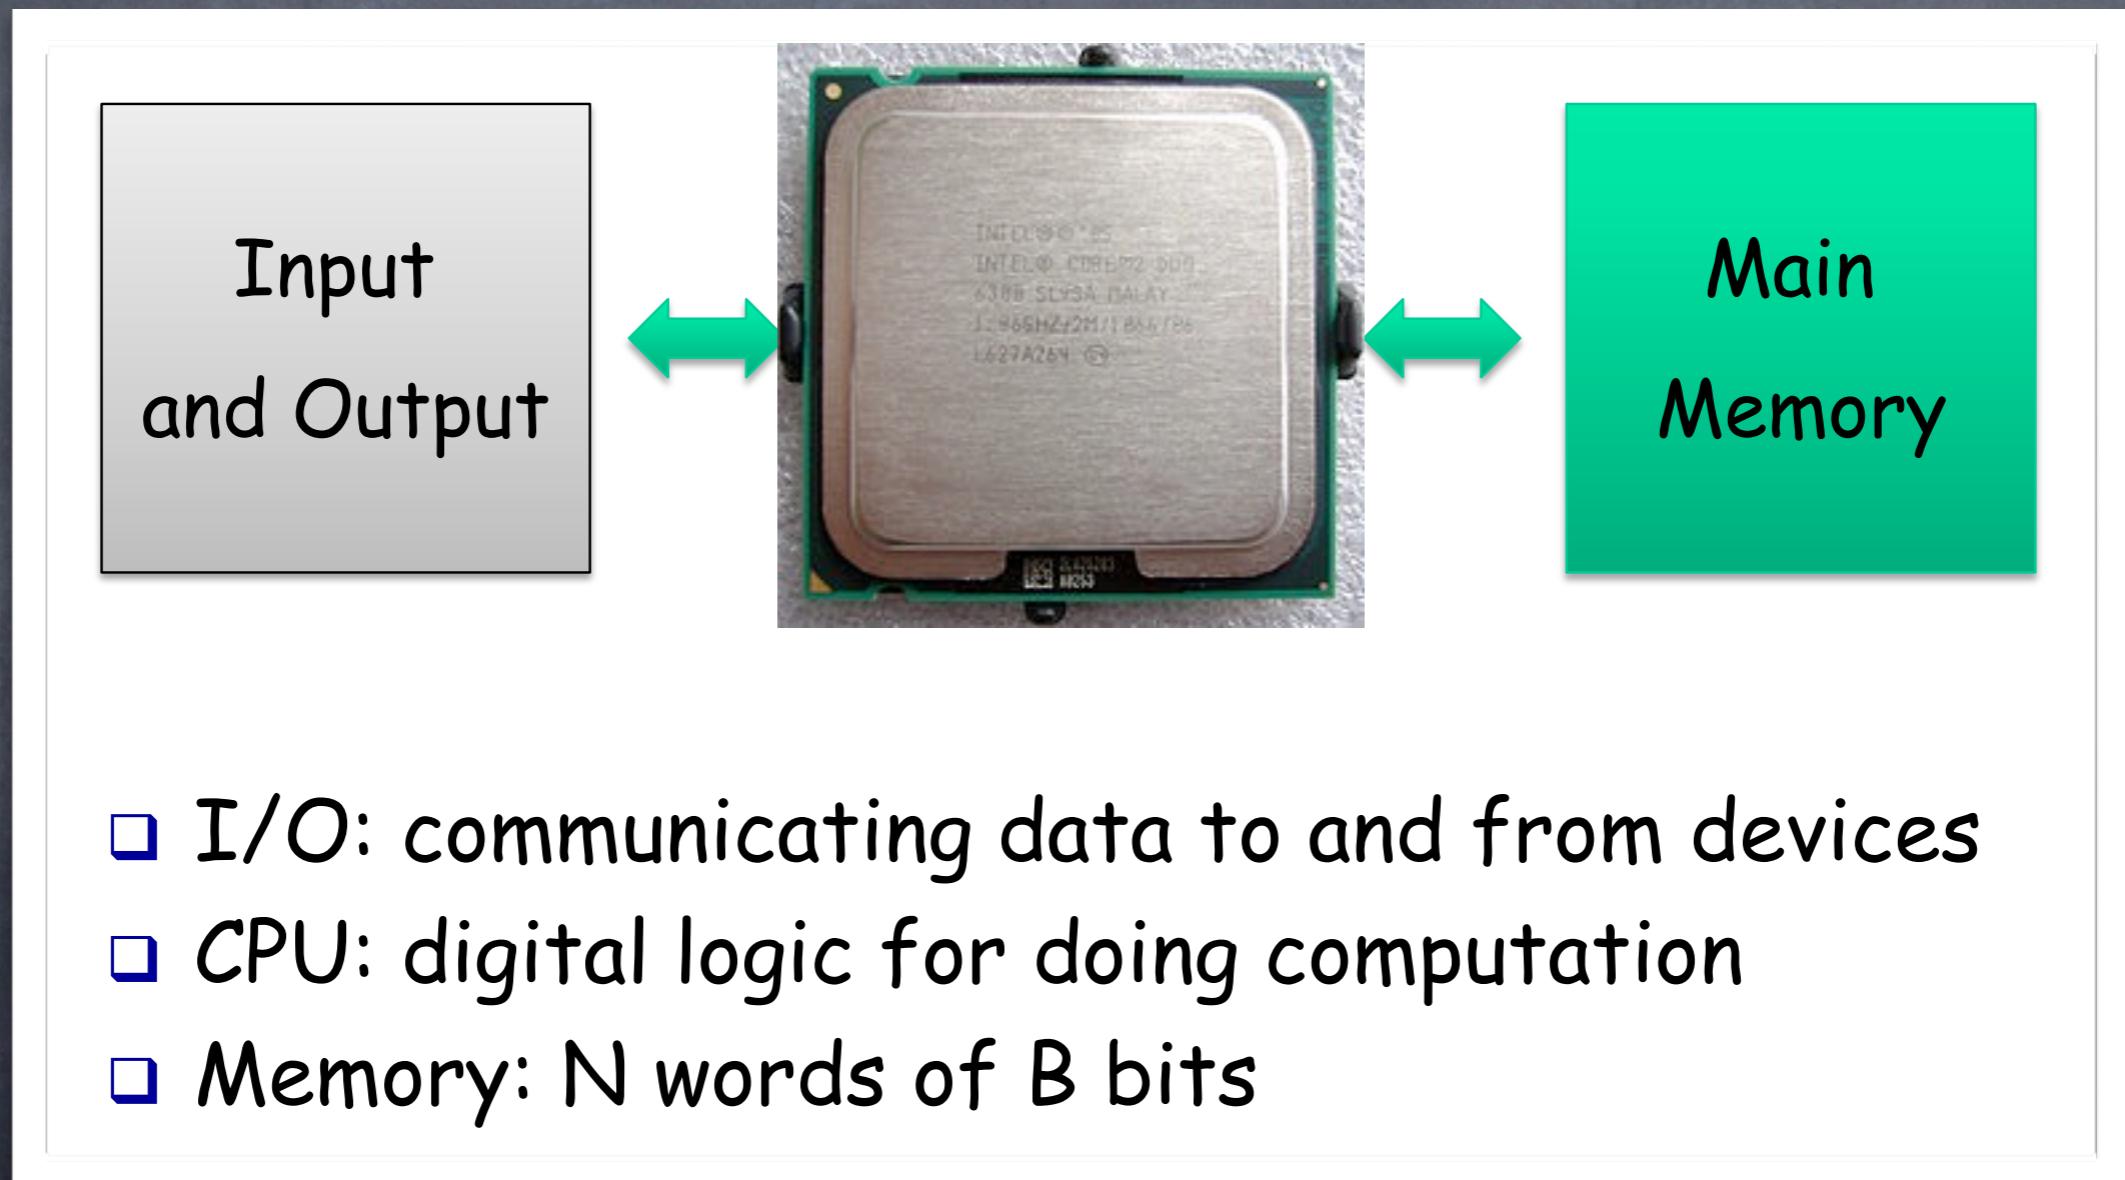
MAC422 Sistemas Operacionais

Prof. Marcel P. Jackowski

Aula #2

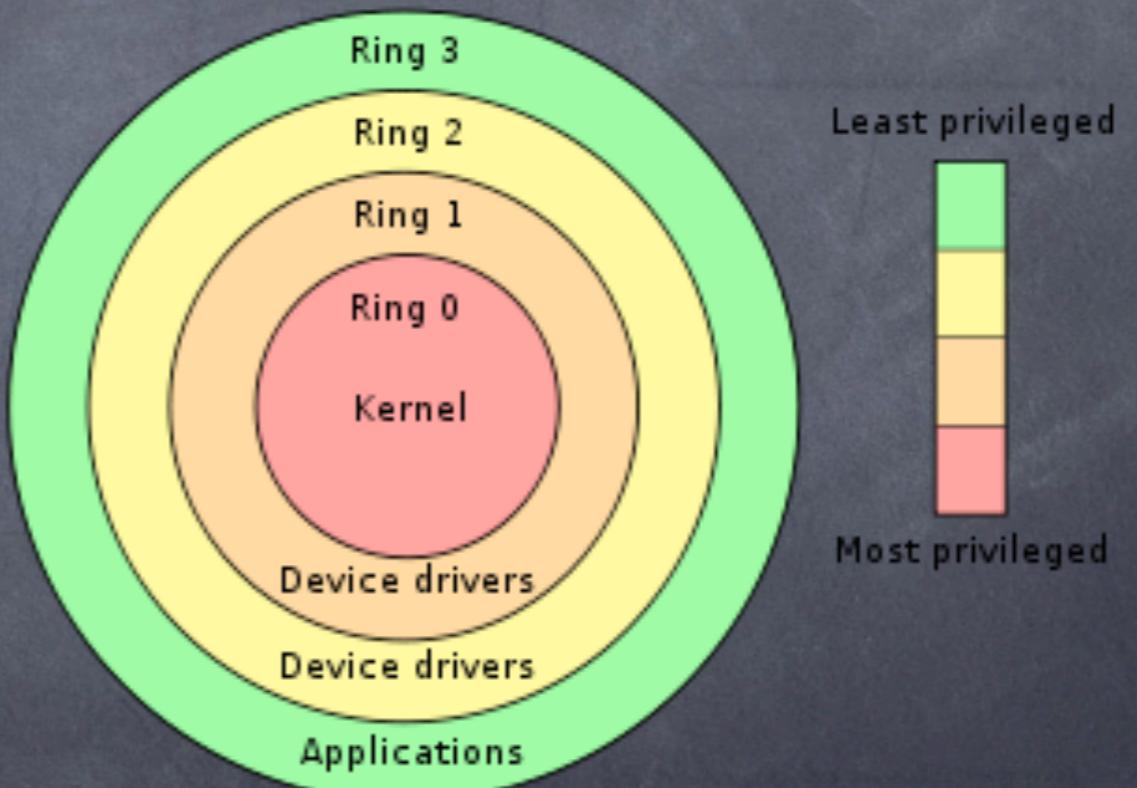
Hardware e Programação x86

Arquitetura de Von Neumann

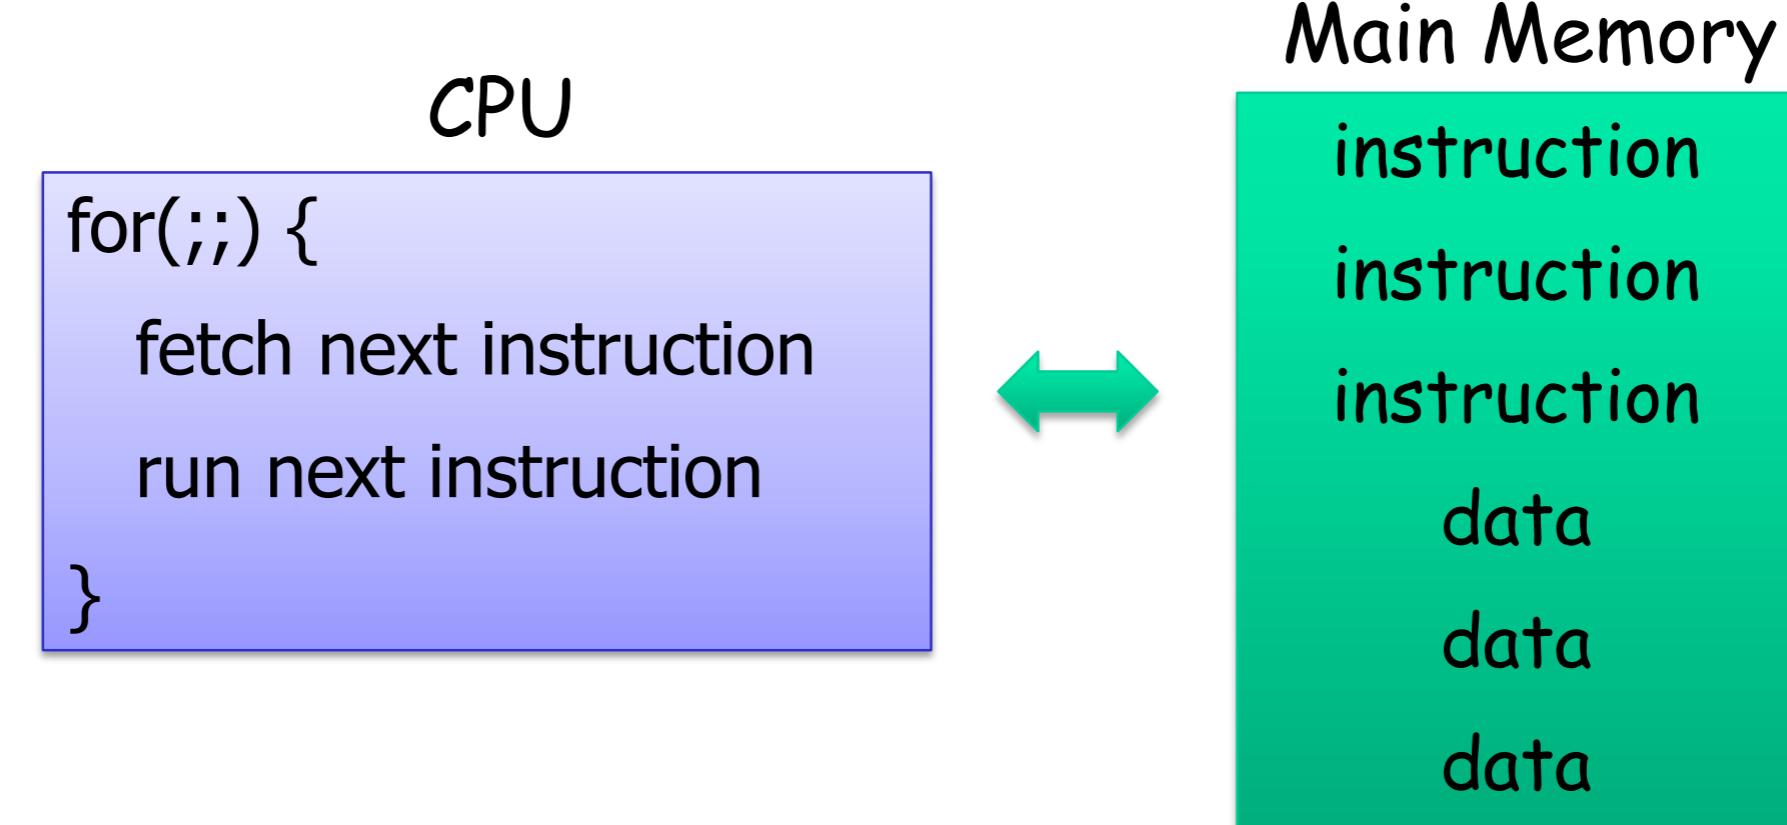


Modos de execução

- Privilegiado
 - Monitor, kernel, supervisor, ring 0 ou sistema.
- Não-privilegiado
 - Modo usuário ou aplicação.

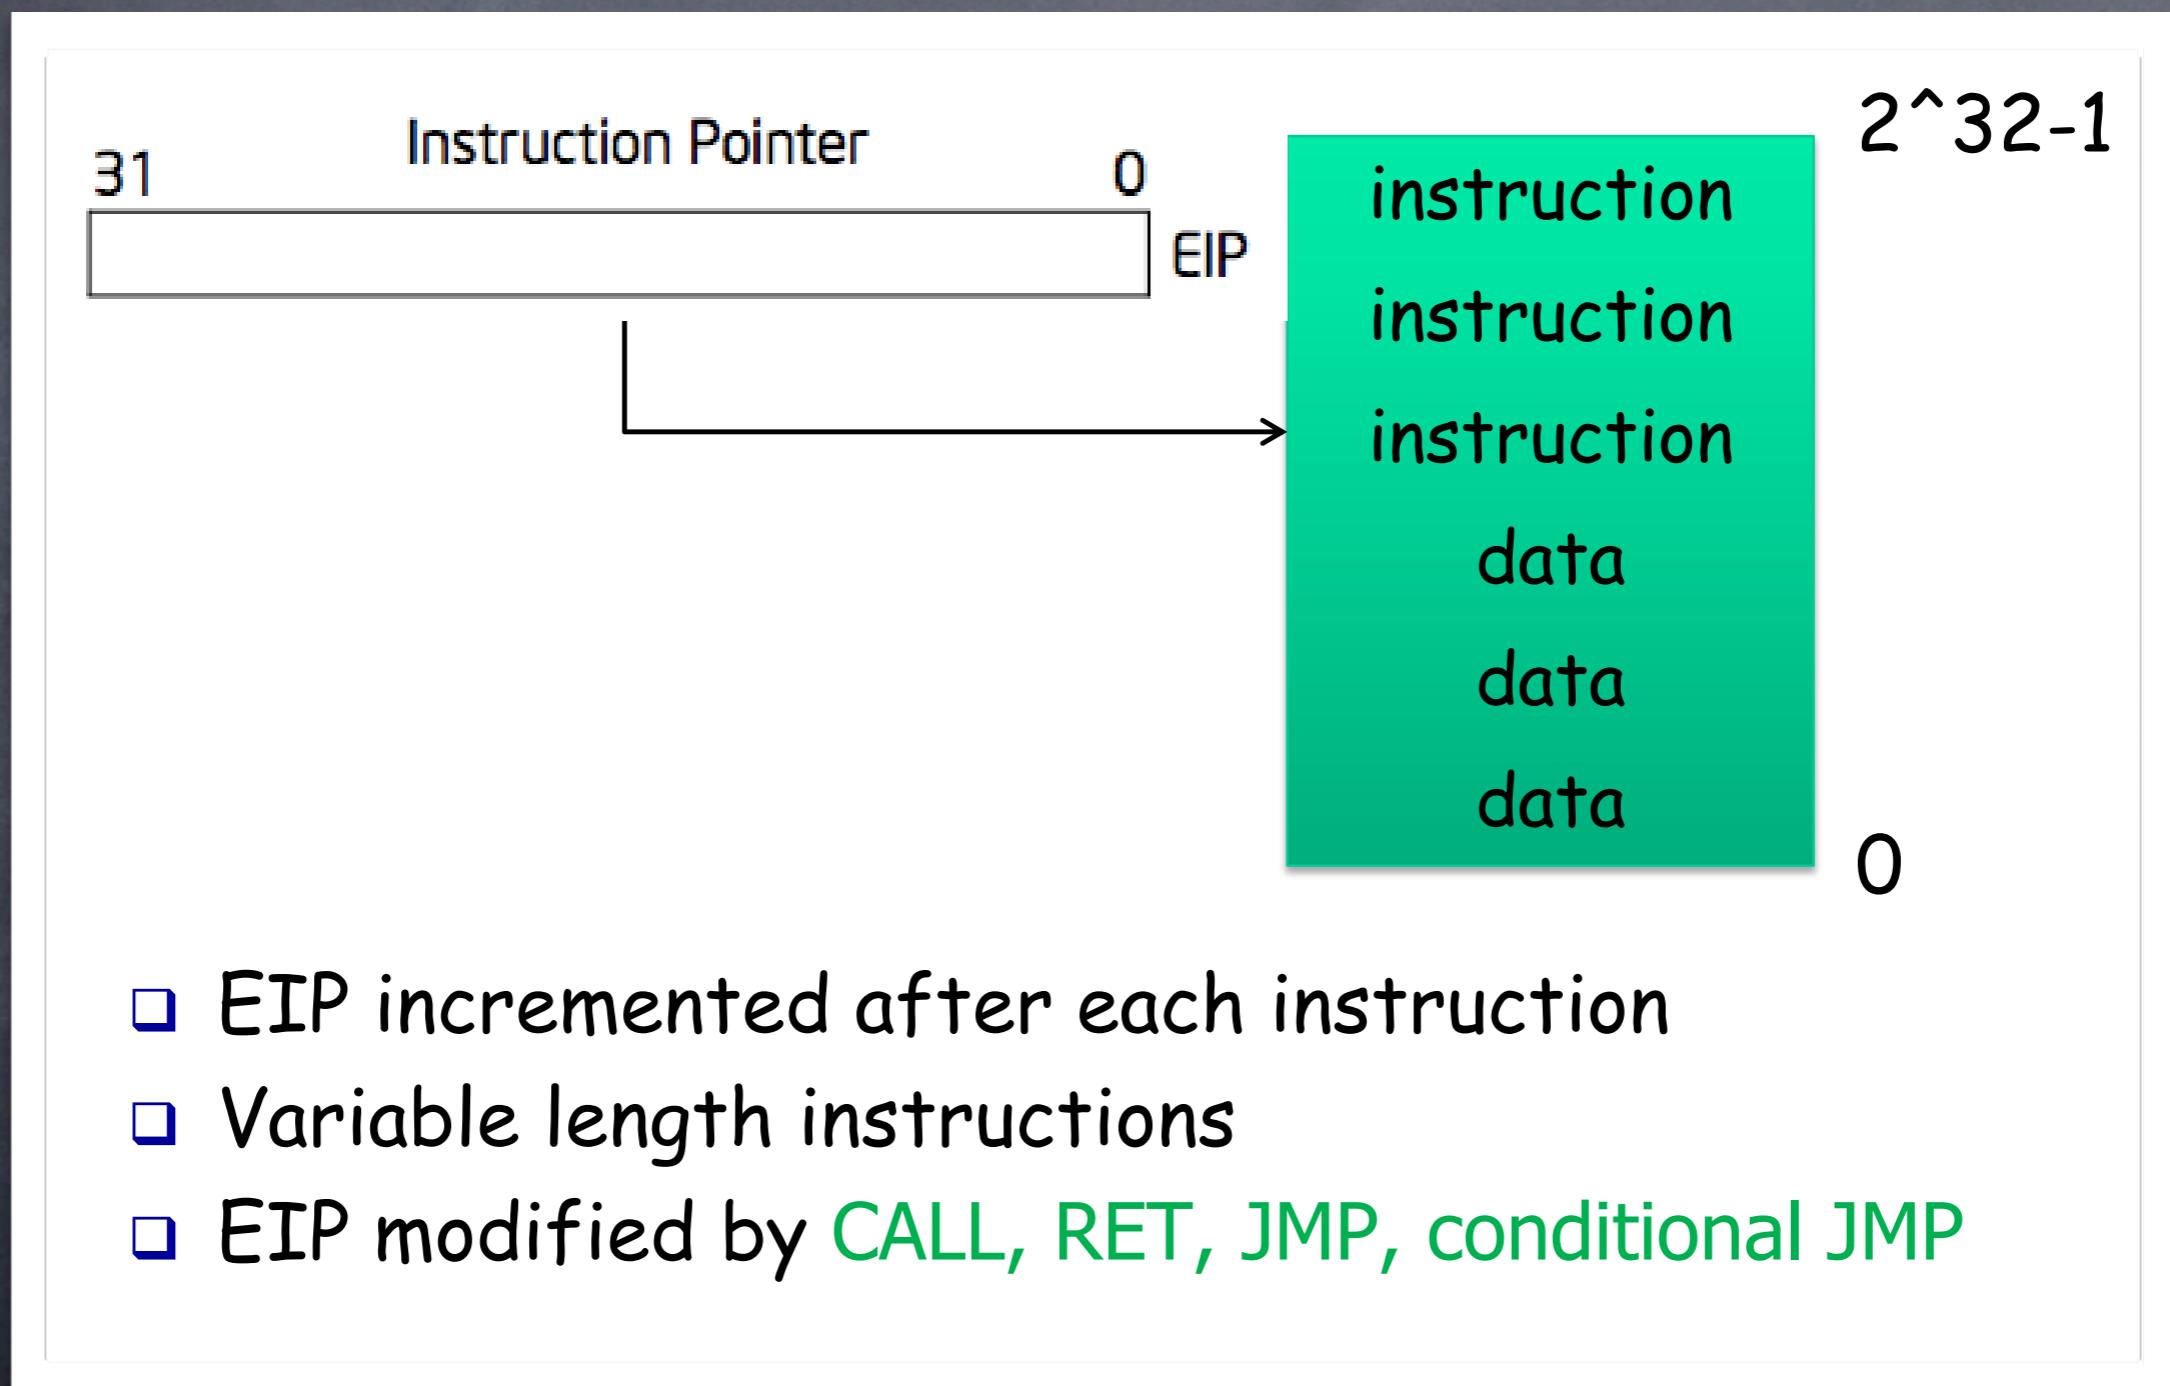


Execução de um programa



- Memory holds both *instructions* and *data*
- CPU interprets instructions
- Instructions read/write data

Contador de programa no x86



E em arquiteturas 64 bits ?

Registradores: espaço de trabalho da CPU

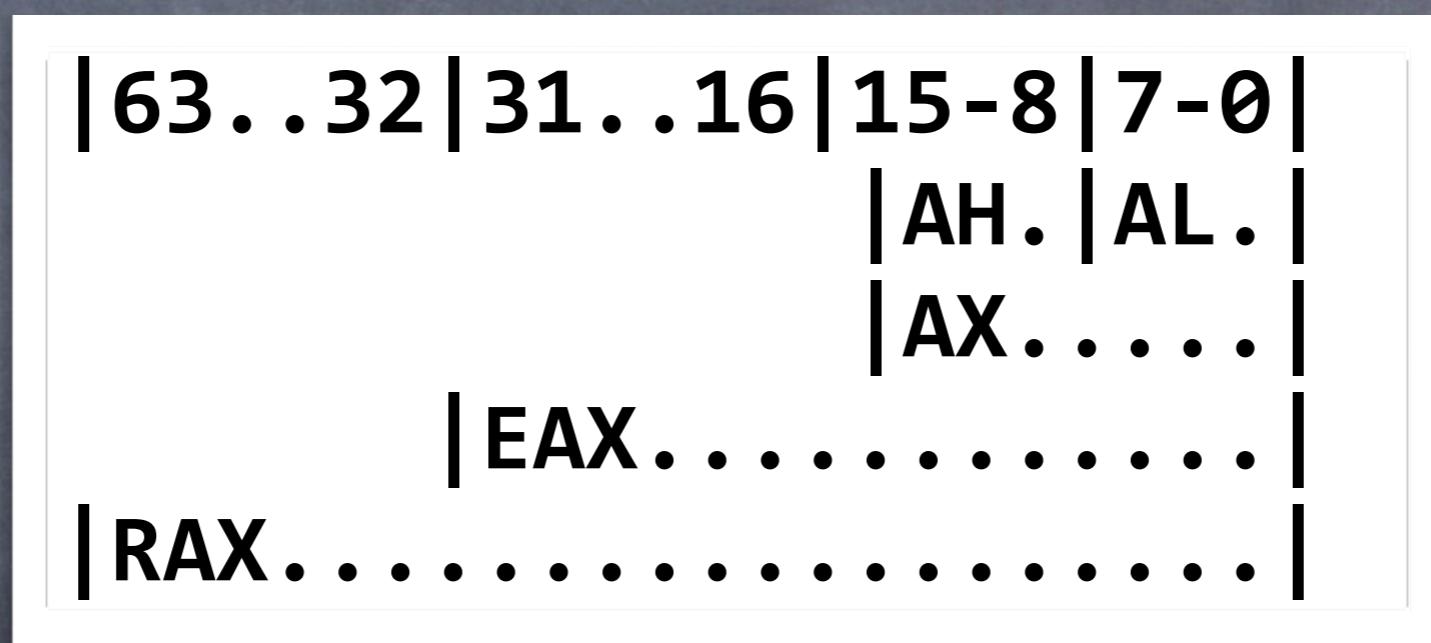
General-Purpose Registers				16-bit	32-bit
31	16 15	8 7	0		
		AH	AL	AX	EAX
		BH	BL	BX	EBX
		CH	CL	CX	ECX
		DH	DL	DX	EDX
		BP		EBP	
		SI		ESI	
		DI		EDI	
		SP		ESP	

- 8, 16, and 32 bit versions
- Example: ADD EAX, 10
 - More: SUB, AND, etc
- By convention some for special purposes

ESP: stack pointer
EBP: frame base pointer
ESI: source index
EDI: destination index

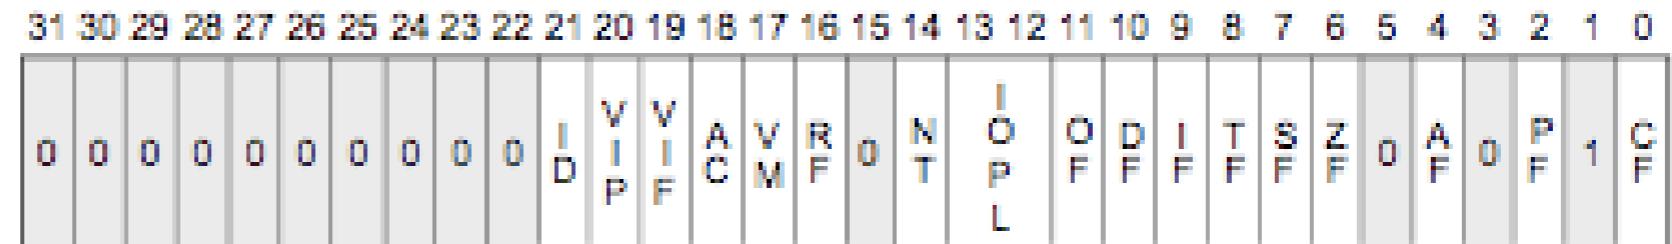
Agora 64 bits

- Exemplo:
 - Acumulador



- Registradores extras:
 - R8-R15: 64 bits

EFLAGS



X ID Flag (ID)

X Virtual Interrupt Pending (VIP)

X Virtual Interrupt Flag (VIF)

X Alignment Check (AC)

X Virtual-8086 Mode (VM)

X Resume Flag (RF)

X Nested Task (NT)

X I/O Privilege Level (IOPL)

S Overflow Flag (OF)

C Direction Flag (DF)

X Interrupt Enable Flag (IF)

X Trap Flag (TF)

S Sign Flag (SF)

S Zero Flag (ZF)

S Auxiliary Carry Flag (AF)

S Parity Flag (PF)

S Carry Flag (CF)

S Indicates a Status Flag

C Indicates a Control Flag

X Indicates a System Flag

- Track current CPU status

TEST EAX, 0
JNZ address

Categorias

□ Instruction classes

- Data movement: MOV, PUSH, POP, ...
- Arithmetic: TEST, SHL, ADD, AND, ...
- I/O: IN, OUT, ...
- Control: JMP, JZ, JNZ, CALL, RET
- String: MOVS, REP, ...
- System: INT, IRET

□ Instruction syntax

- Intel manual Volumne 2: op dst, src
- AT&T (gcc/gas): op src, dst
 - op uses suffix b, w, l for 8, 16, 32-bit operands

Sintaxe Intel vs. AT&T

+-----+-----+	
Intel Code AT&T Code	
+-----+-----+	
mov eax, 1	movl \$1,%eax
mov ebx, 0ffh	movl \$0xff,%ebx
int 80h	int \$0x80
mov ebx, eax	movl %eax,%ebx
mov eax, [ecx]	movl(%ecx),%eax
mov eax, [ebx+3]	movl 3(%ebx),%eax
mov eax, [ebx+20h]	movl 0x20(%ebx),%eax
add eax, [ebx+ecx*2h]	addl(%ebx,%ecx,0x2),%eax
lea eax, [ebx+ecx]	leal(%ebx,%ecx),%eax
sub eax, [ebx+ecx*4h-20h]	subl -0x20(%ebx,%ecx,0x4),%eax

Assembly inline (gcc)

- Can embed assembly code in C code
 - Many examples in xv6
- Basic syntax: `asm ("assembly code")`
e.g., `asm ("movl %eax %ebx")`
- Advanced syntax:
`asm (assembler template`
 `: output operands /* optional */`
 `: input operands /* optional */`
 `: list of clobbered registers /* optional */);`
e.g., `int val;`
`asm ("movl %%ebp,%0" : "=r" (val));`

Mais exemplos

```
int a=10, b;
asm ("movl %1, %%eax;
      movl %%eax, %0;"
      : "=r"(b)          /* output */
      : "r"(a)           /* input */
      : "%eax"           /* clobbered register */
);
asm ("movl %0,%%eax;
      movl %1,%%ecx;
      call _foo"
      : /* no outputs */
      : "g" (from), "g" (to)
      : "eax", "ecx"
);
```

Exemplo de uso dos EFLAGS

```
int main()
{
    int x = 0;
    if (x < 0)
        x++;
    else
        x--;
    return 0;
}
```

```
gcc -c -S -m32 if.c
```

```
.text
.globl _main
_main:
    pushl    %ebp
    movl    %esp, %ebp
    subl    $24, %esp
    movl    $0, -12(%ebp)
    cmpl    $0, -12(%ebp)
    jns     L2
    incl    -12(%ebp)
    jmp     L4

L2:
    decl    -12(%ebp)

L4:
    movl    $0, %eax
    leave
    ret
.subsections_via_symbols
```

RAM: maior espaço de trabalho

movl %eax, %edx	edx = eax;	<i>register mode</i>
movl \$0x123, %edx	edx = 0x123;	<i>immediate</i>
movl 0x123, %edx	edx = *(int32_t*)0x123;	<i>direct</i>
movl (%ebx), %edx	edx = *(int32_t*)ebx;	<i>indirect</i>
movl 4(%ebx), %edx	edx = *(int32_t*)(ebx+4);	<i>displaced</i>

- Memory instructions: **MOV, PUSH, POP, etc**
- Most instructions can take a memory address

Pilha: estrutura de dados em RAM

Example instruction What it does

pushl %eax

subl \$4, %esp
movl %eax, (%esp)

popl %eax

movl (%esp), %eax
addl \$4, %esp

call 0x12345

pushl %eip (*)
movl \$0x12345, %eip (*)

ret

popl %eip (*)

- For implementing function calls
- Stack grows “down” on x86

Alguns detalhes...

- 8086 16-bit register and 20-bit bus addresses
- These extra 4 bits come from *segment register*
 - CS: code segment, for EIP
 - Instruction address: CS * 16 + EIP
 - SS: stack segment, for ESP and EBP
 - DS: data segment for load/store via other registers
 - ES: another data segment, destination for string ops
- Make life more complicated
 - Cannot directly use 16-bit stack address as pointer
 - For a far pointer programmer must specify segment reg
 - Pointer arithmetic and array indexing across seg bound

Mais alguns detalhes...

- 80386: 32 bit register and addresses (1985)
- AMD k8: 64 bit (2003)
 - RAX instead of EAX
 - x86-64, x64, amd64, intel64: all same thing
- Backward compatibility
 - Boots in 16-bit mode; bootasm.S switches to 32
 - Prefix 0x66 gets 32-bit mode instructions
 - MOVW in 32-bit mode = 0x66 + MOVW in 16-bit mode
 - .code32 in bootasm.S tells assembler to insert 0x66
- 80386 also added virtual memory addresses

Entrada/Saída

```
#define DATA_PORT      0x378
#define STATUS_PORT    0x379
#define BUSY 0x80
#define CONTROL_PORT 0x37A
#define STROBE 0x01
void
lpt_putc(int c)
{
    /* wait for printer to consume previous byte */
    while((inb(STATUS_PORT) & BUSY) == 0)
        ;

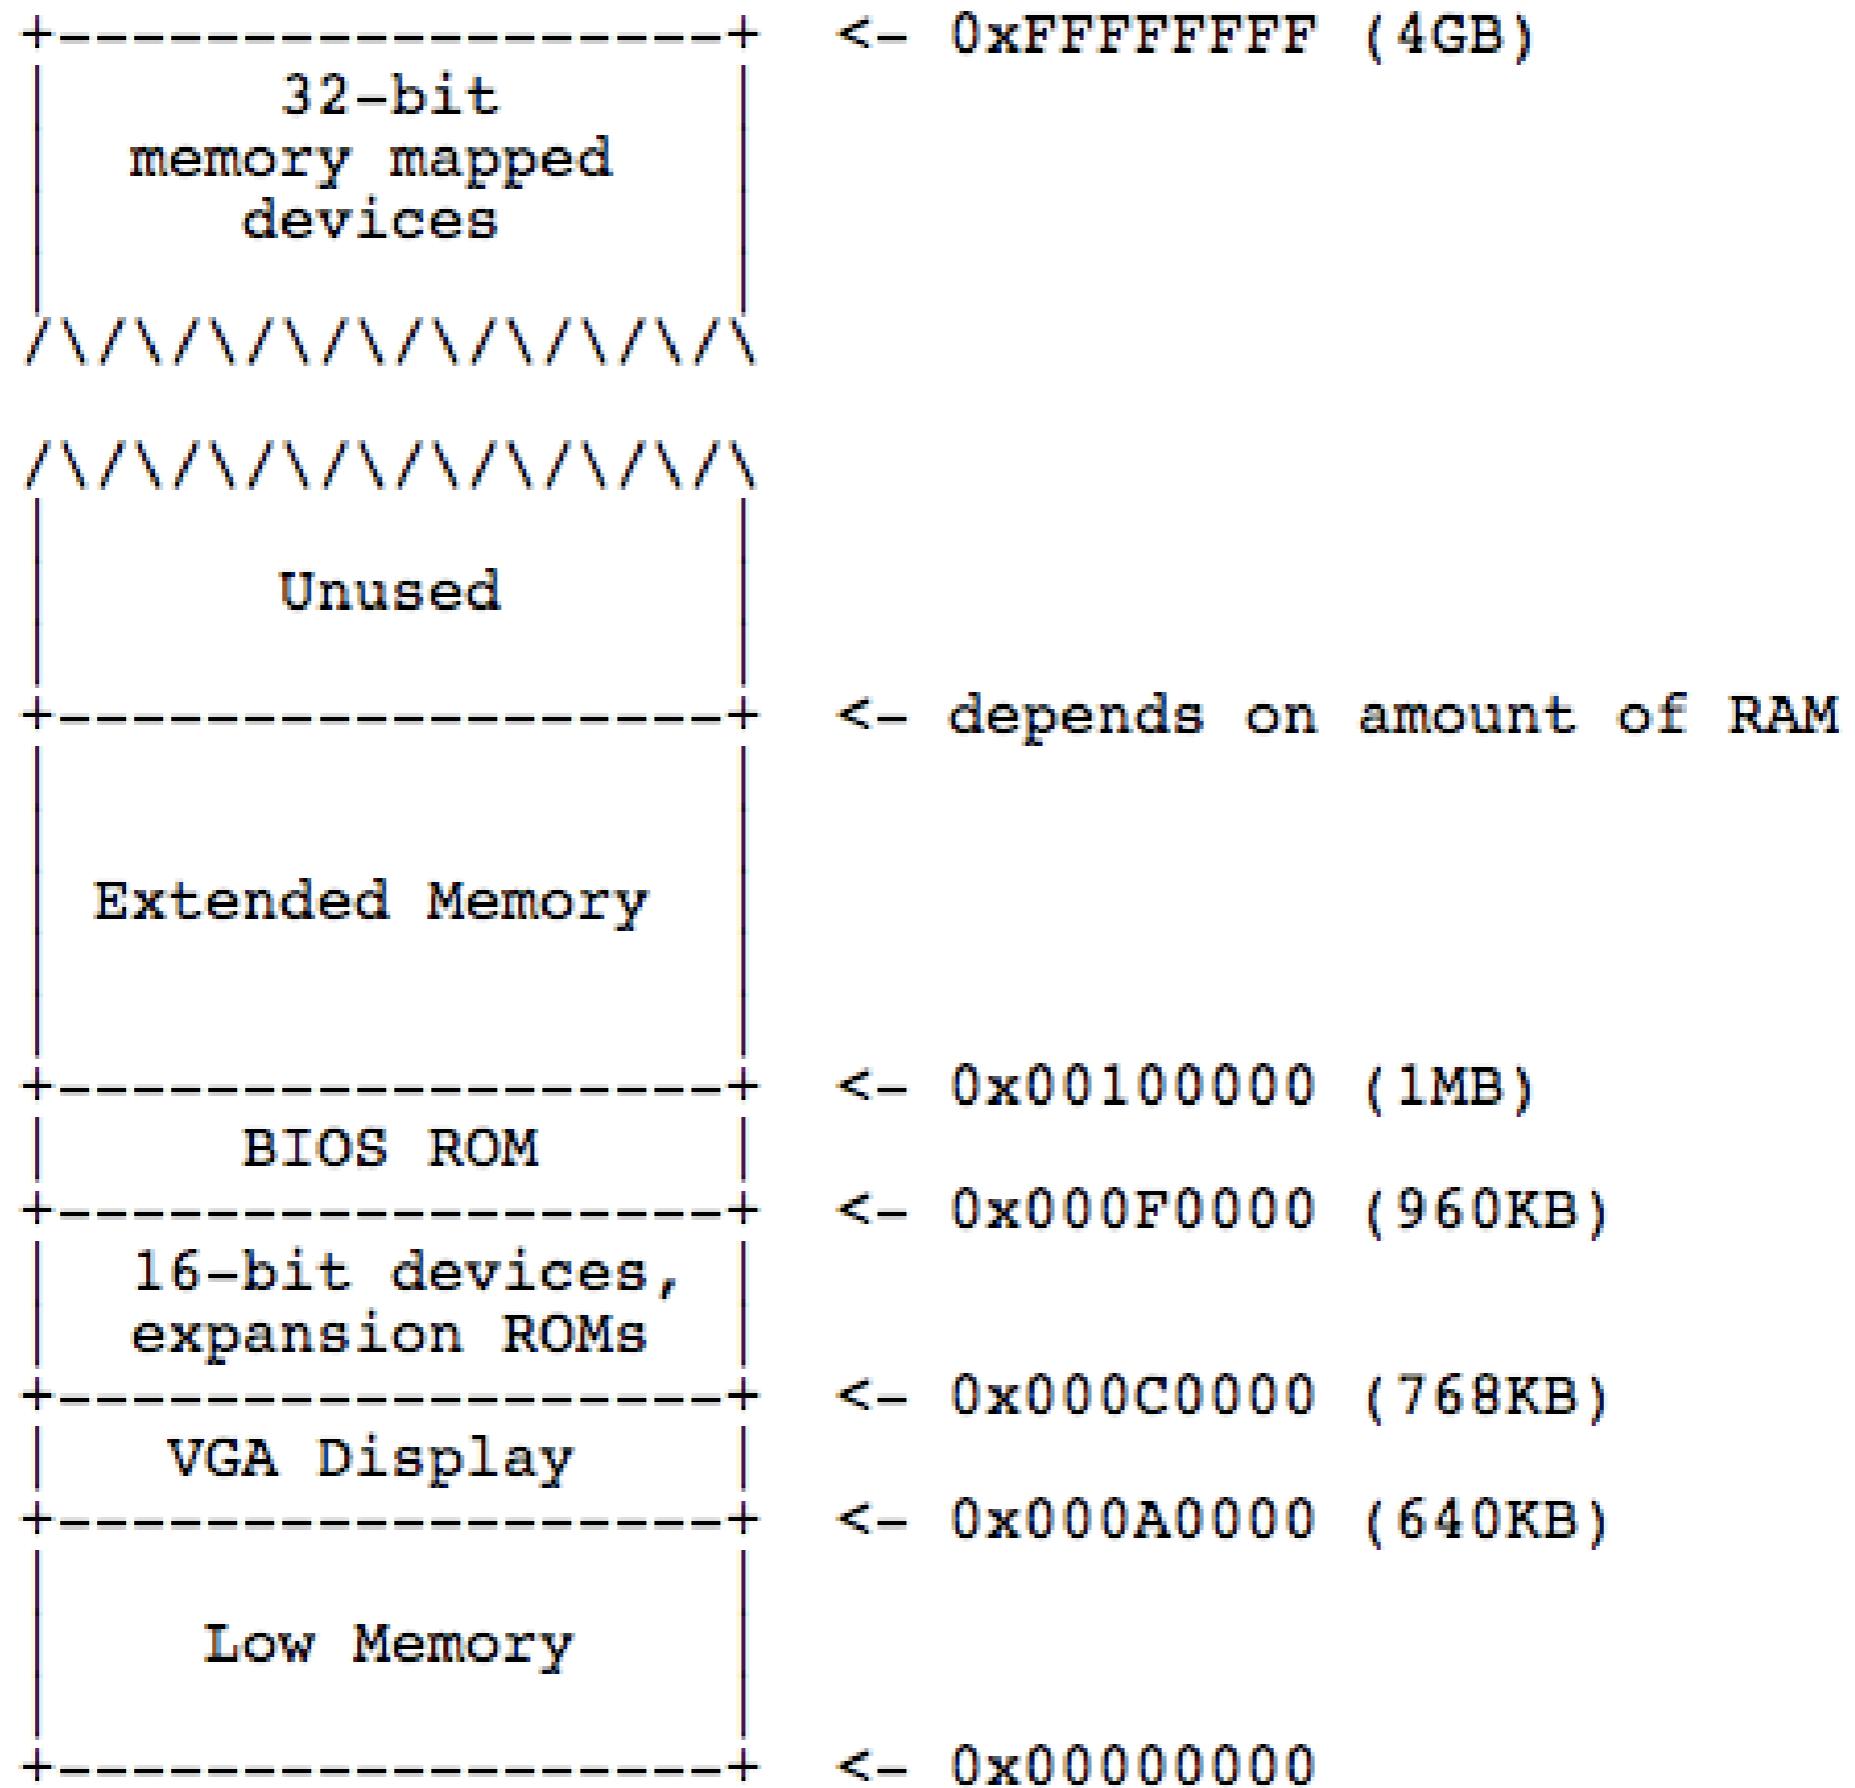
    /* put the byte on the parallel lines */
    outb(DATA_PORT, c);

    /* tell the printer to look at the data */
    outb(CONTROL_PORT, STROBE);
    outb(CONTROL_PORT, 0);
}
```

- 8086: only 1024 addresses

E/S mapeada em memória

- Use normal addresses for I/O
 - No special instructions
 - No 1024 limit
 - Hardware routes to device
- Works like “magic” memory
 - I/O device addressed and accessed like memory
 - However, reads and writes have “side effects”
 - Read result can change due to external events



Convenções do gcc (i): chamadas de funções

```
1 subprogram_label:  
2     push    ebp          ; save original EBP value on stack  
3     mov     ebp, esp    ; new EBP = ESP  
4 ; subprogram code  
5     pop    ebp          ; restore original EBP value  
6     ret
```

Figure 4.3: General subprogram form

```
1     push    dword 1      ; pass 1 as parameter  
2     call    fun  
3     add    esp, 4       ; remove parameter from stack
```

Figure 4.5: Sample subprogram call

```

main() {
    return foo(10, 20);
}

int foo(int x, inty) {
    return x+y;
}

```

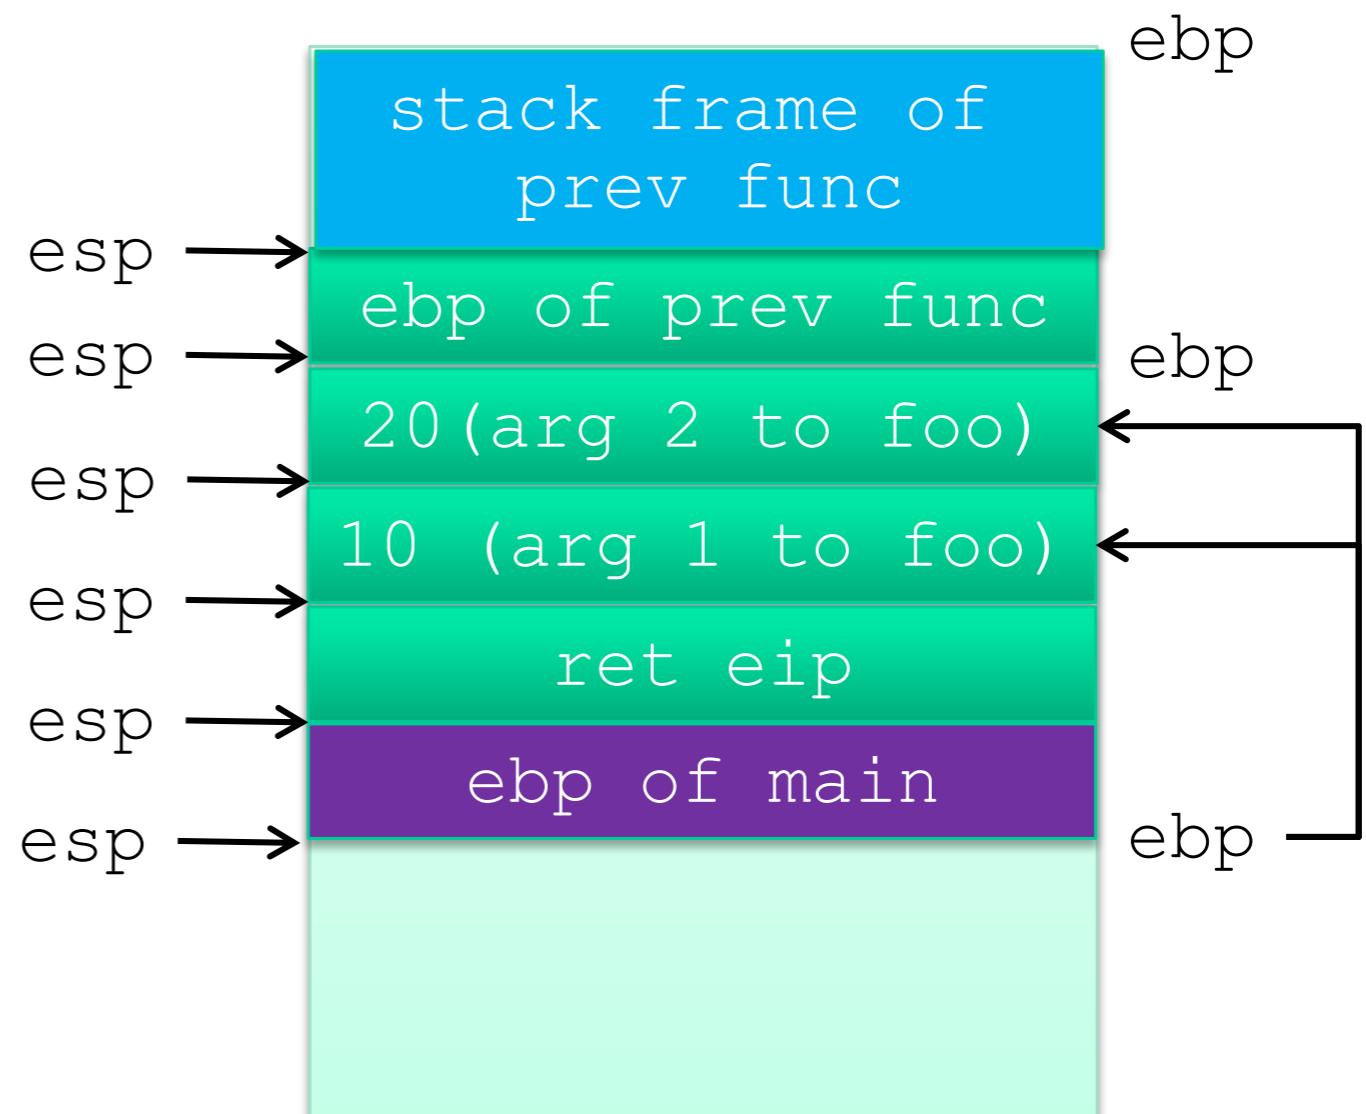
```

_main:
→ pushl %ebp
→ movl %esp, %ebp
→ pushl $20
→ pushl $10
→ call foo
→ movl %ebp, %esp
→ popl %ebp
→ ret

_foo:
→ pushl %ebp
→ movl %esp, %ebp
→ movl 0xc(%ebp), %eax
→ add 0x8(%ebp), %eax
→ movl %ebp, %esp
→ popl %ebp
→ ret

```

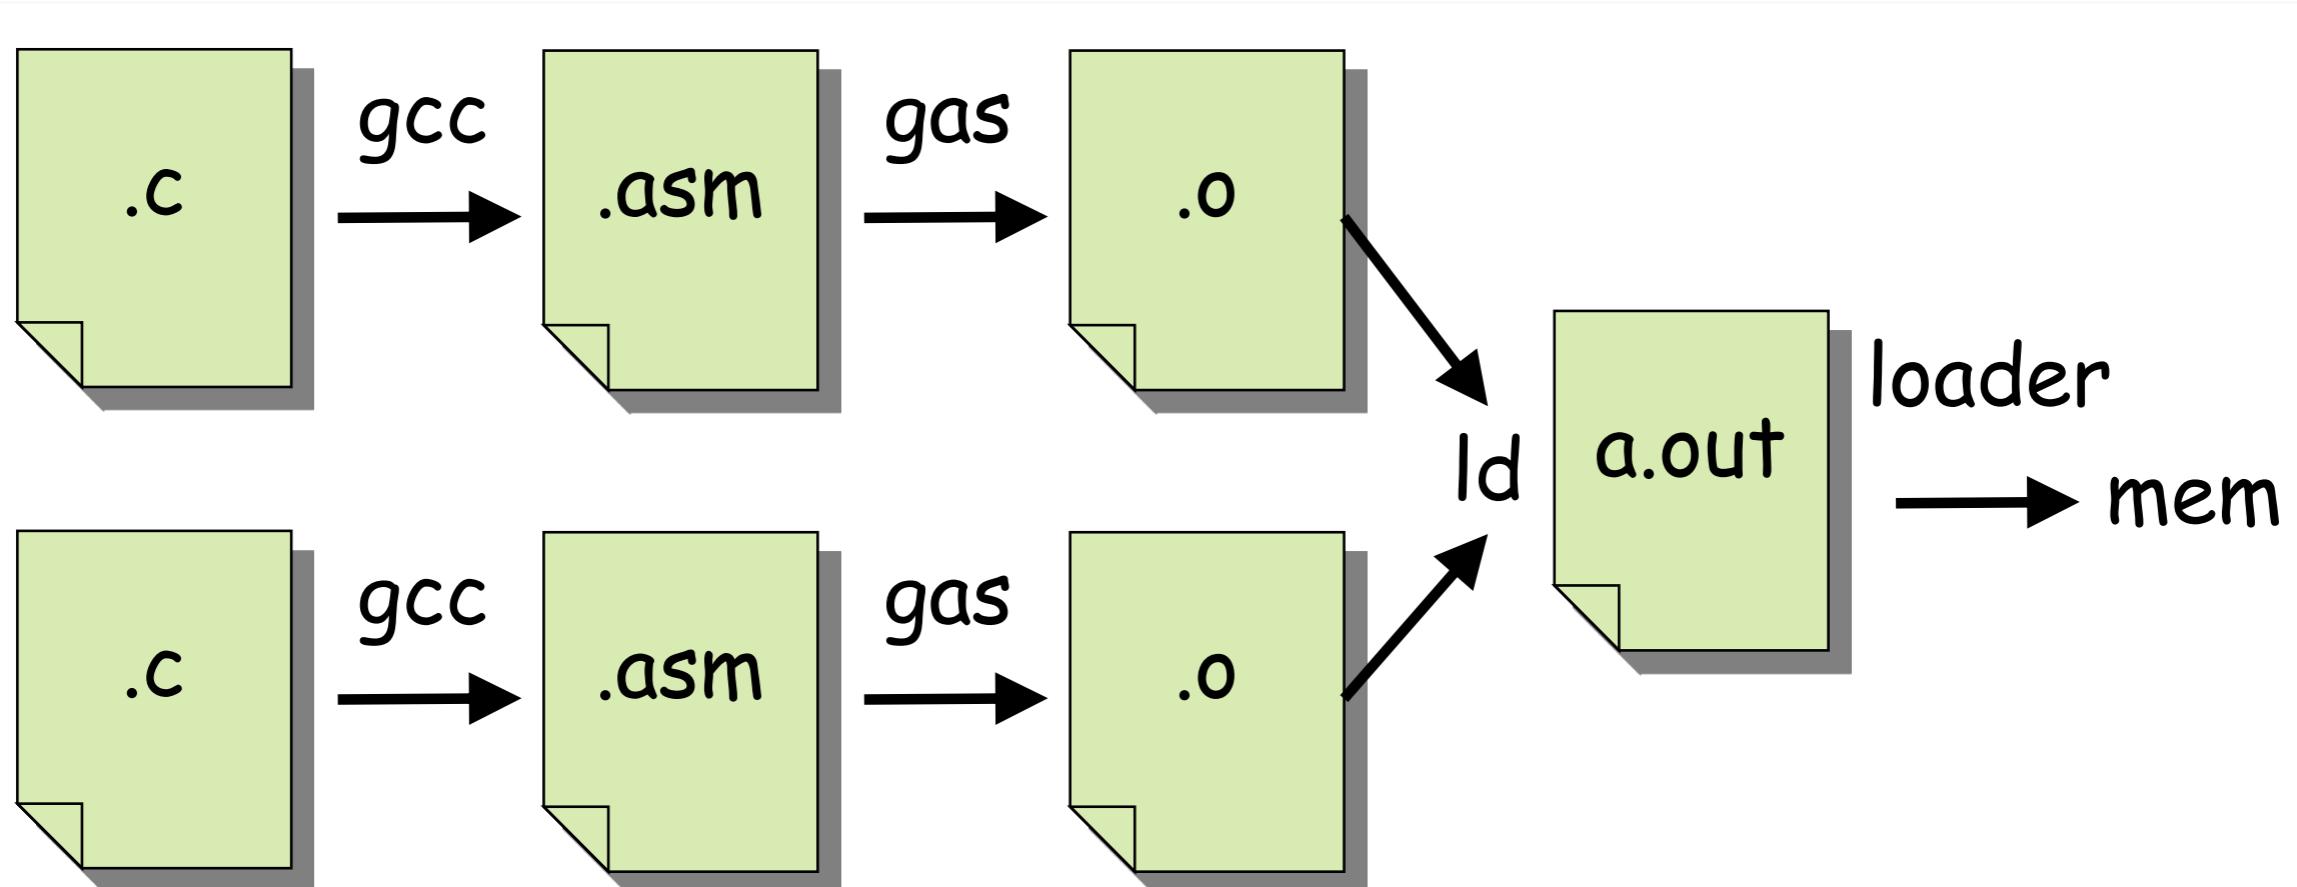
Exemplo



Convenções do gcc (ii): registradores

- %eax contains return value, %ecx, %edx may be trashed
 - 64 bit return value: %eax + %edx
- %ebp, %ebx, %esi, %edi must be as before call
- Caller saved: %eax, %ecx, %edx
- Callee saved: %ebp, %ebx, %esi, %edi

Do código-fonte até o executável...



- Compiler, assembler, linker, and loader

Emulador

- QEMU pc emulator
 - Does what a real PC does
 - Except implemented in s/w!

- Run like a normal program
on “host” OS

xv6

PC Emulator

Linux

PC

Emulador de Registradores

```
int32_t regs[8];
#define REG_EAX 1;
#define REG_EBX 2;
#define REG_ECX 3;
...
int32_t eip;
int16_t segreggs[4];
...
```

Emulador de lógica da CPU

```
for (;;) {
    read_instruction();
    switch (decode_instruction_opcode()) {
        case OPCODE_ADD:
            int src = decode_src_reg();
            int dst = decode_dst_reg();
            regs[dst] = regs[dst] + regs[src];
            break;
        case OPCODE_SUB:
            int src = decode_src_reg();
            int dst = decode_dst_reg();
            regs[dst] = regs[dst] - regs[src];
            break;
        ...
    }
    eip += instruction_length;
}
```

Emulador de memória

```
uint8_t read_byte(uint32_t phys_addr) {
    if (phys_addr < LOW_MEMORY)
        return low_mem[phys_addr];
    else if (phys_addr >= 960*KB && phys_addr < 1*MB)
        return rom_bios[phys_addr - 960*KB];
    else if (phys_addr >= 1*MB && phys_addr < 1*MB+EXT_MEMORY) {
        return ext_mem[phys_addr-1*MB];
    } else ...
}

void write_byte(uint32_t phys_addr, uint8_t val) {
    if (phys_addr < LOW_MEMORY)
        low_mem[phys_addr] = val;
    else if (phys_addr >= 960*KB && phys_addr < 1*MB)
        ; /* ignore attempted write to ROM! */
    else if (phys_addr >= 1*MB && phys_addr < 1*MB+EXT_MEMORY) {
        ext_mem[phys_addr-1*MB] = val;
    } else ...
}
```