



MAC422 Sistemas Operacionais

Prof. Marcel P. Jackowski

Aula #4

Criação de processos e mecanismos de comunicação

Criação de processos (i)

- Em Unix, Linux e variações:
 - Utiliza-se a função `fork()` para criar um novo processo e `exec()` para executar um programa externo.
 - Por que não uma única chamada `fork_and_exec()` ?
 - Pense como você implementaria “ls | wc” ?



Criação de processos (ii)

- Option 1: **cloning** (e.g., Unix `fork()`, `exec()`)
 - Pause current process and save its state
 - Copy its PCB (can select what to copy)
 - Add new PCB to ready queue
 - Must distinguish parent and child
- Option 2: **from scratch** (Win32 `CreateProcess`)
 - Load code and data into memory
 - Create and initialize PCB (make it like saved from context switch)
 - Add new PCB to ready queue

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int main()
{
    int a = 2;
    int b = -1;
    int c = 3;

    fork();
    a = a + 2;

    printf("pid=%d, a=%d, b=%d, c=%d\n",getpid(),a,b,c);

    fork();

    b = a + c;
    printf("pid=%d, a=%d, b=%d, c=%d\n",getpid(),a,b,c);

    fork();

    c = a + b;
    printf("pid=%d, a=%d, b=%d, c=%d\n",getpid(),a,b,c);

    exit(0);
}
```

- Quantos processos serão criados ?
- Quantas vezes serão impressos os valores de a, b e c ?
- Eles serão diferentes entre execuções ?

Término de processos

- Normal: `exit(int status)`
 - OS passes exit status to parent via `wait(int *status)`
 - OS frees process resources
- Abnormal: `kill(pid_t pid, int sig)`
 - OS can kill process
 - Process can kill process

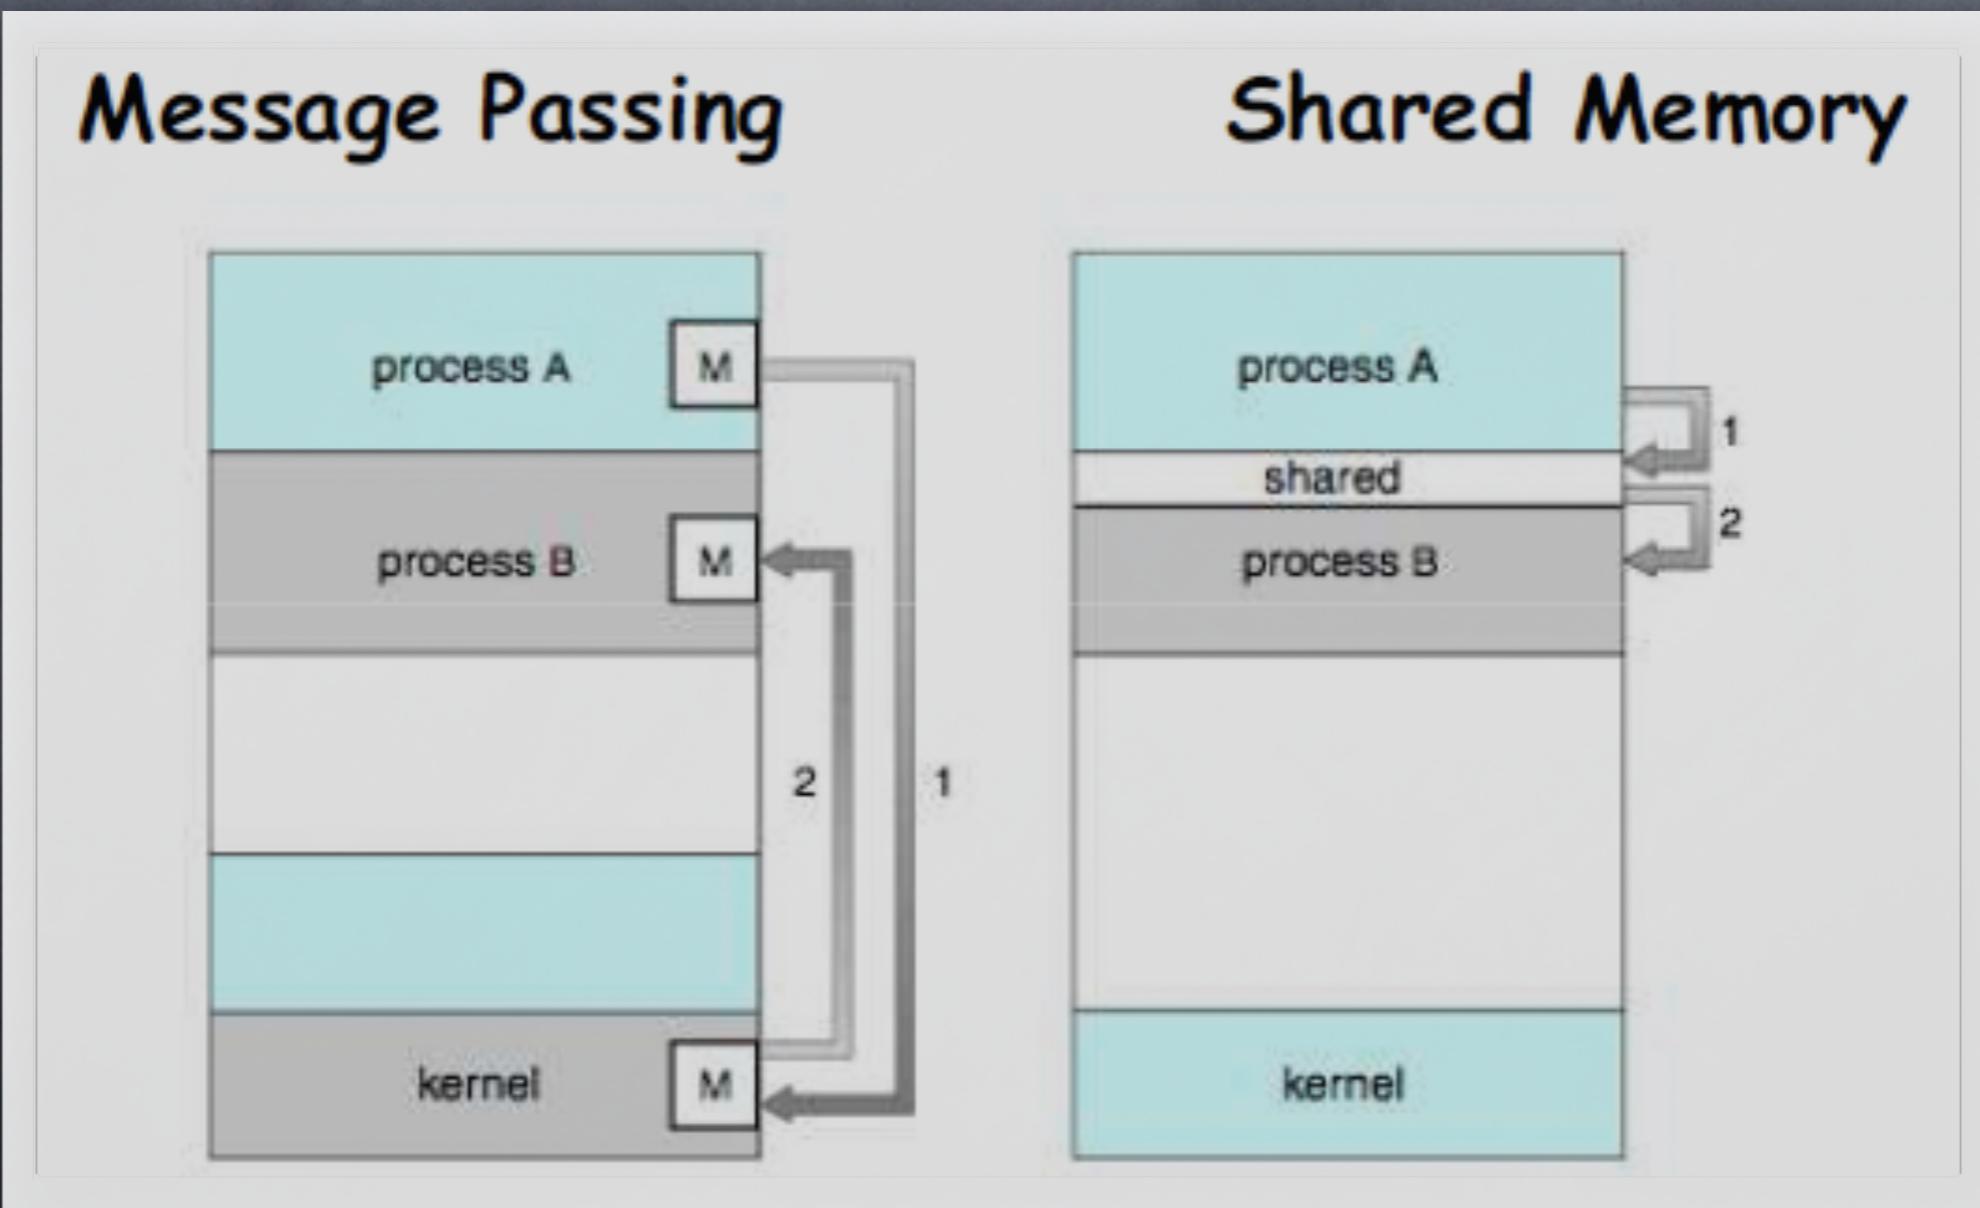
Entre zumbis e órfãos

- What if child exits before parent?
 - Child becomes zombie
 - Need to store exit status
 - OS can't fully free
 - Parent must call wait() to reap child
- What if parent exits before child?
 - Child becomes orphan
 - Need some process to query exit status and maintain process tree
 - Re-parent to the first process, the init process

Processos cooperativos

- **Independent** process cannot affect or be affected by the execution of another process.
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity/Convenience

Comunicação interprocessos (IPC)



Troca de mensagens

Memória compartilhada

Vantagens e desvantagens

□ Message passing

- Why good?
- Why bad?

□ Shared Memory

- Why good?
- Why bad?

Exemplo de IPC: Sinais em Unix/Linux

- Signals
 - A very short message: just a small integer
 - A fixed set of available signals. Examples:
 - 9: kill
 - 11: segmentation fault
- Installing a handler for a signal
 - `sighandler_t signal(int signum, sighandler_t handler);`
- Send a signal to a process
 - `kill(pid_t pid, int sig)`

```
#include <signal.h>
#include <stdio.h>
#include <stdlib.h>

void sigproc()
{ /* NOTE some versions of UNIX will reset signal to default
   after each call. So for portability reset signal each time
   signal(SIGINT, sigproc);
 */
    printf("you have pressed ctrl-c...\n");
}

void sigkill()
{
    printf("a-ha! you can't kill me!\n");
}

void sigusr1()
{
    printf("Hi, what do you want?\n");
}

int main()
{
    signal(SIGINT, sigproc);
    signal(SIGKILL, sigkill);
    signal(SIGUSR1, sigusr1);
    for(;;) /* infinite loop */
}
```

Exemplo de IPC: Pipes

□ `int pipe(int fds[2])`

- Creates a one way communication channel
- `fds[2]` is used to return two file descriptors
- Bytes written to `fds[1]` will be read from `fds[0]`

```
int pipefd[2];
pipe(pipefd);
switch(pid=fork()) {
case -1: perror("fork"); exit(1);
case 0: close(pipefd[0]);
          // write to fd 1
          break;
default: close(pipefd[1]);
          // read from fd 0
          break;
}
```

Exemplo de IPC: Pipes (ii)

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(void)
{
    int pfds[2];

    pipe(pfds);

    if (!fork()) {
        close(1);          /* close normal stdout */
        dup(pfds[1]);    /* make stdout same as pfds[1] */
        close(pfds[0]); /* we don't need this */
        execlp("ls", "ls", NULL);
    } else {
        close(0);          /* close normal stdin */
        dup(pfds[0]);    /* make stdin same as pfds[0] */
        close(pfds[1]); /* we don't need this */
        execlp("wc", "wc", "-l", NULL);
    }

    return 0;
}
```

ls | wc

Exemplo de IPC: Memória compartilhada

- `int shmget(key_t key, size_t size, int shmflg);`
 - Create a shared memory segment
 - key: unique identifier of a shared memory segment, or `IPC_PRIVATE`
- `int shmat(int shmid, const void *addr, int flg)`
 - Attach shared memory segment to address space of the calling process
 - shmid: id returned by `shmget()`
- `int shmdt(const void *shmaddr);`
 - Detach from shared memory
- Problem: synchronization!

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define SHM_SIZE 1024 /* make it a 1K shared memory segment */

int main(int argc, char *argv[ ])
{
    key_t key;
    int shmid;
    char *data;
    int mode;

    if (argc > 2) {
        fprintf(stderr, "usage: shmdemo [data_to_write]\n");
        exit(1);
    }

    /* make the key: */
    if ((key = ftok("shmdemo.c", 'R')) == -1) {
        perror("ftok");
        exit(1);
    }

    /* connect to (and possibly create) the segment: */
    if ((shmid = shmget(key, SHM_SIZE, 0644 | IPC_CREAT)) == -1) {
```

```
/* connect to (and possibly create) the segment: */
if ((shmid = shmget(key, SHM_SIZE, 0644 | IPC_CREAT)) == -1) {
    perror("shmget");
    exit(1);
}

/* attach to the segment to get a pointer to it: */
data = shmat(shmid, (void *)0, 0);
if (data == (char *)(-1)) {
    perror("shmat");
    exit(1);
}

/* read or modify the segment, based on the command line: */
if (argc == 2) {
    printf("writing to segment: \"%s\"\n", argv[1]);
    strncpy(data, argv[1], SHM_SIZE);
} else
    printf("segment contains: \"%s\"\n", data);

/* detach from the segment: */
if (shmdt(data) == -1) {
    perror("shmdt");
    exit(1);
}

return 0;
}
```

Informações sobre processos

- ❑ `ps`
- ❑ `top`
- ❑ For each process, there is a corresponding directory `/proc/<pid>` to store this process information in the `/proc` pseudo file system

ls /proc

```
mjack@ubuntu: ~/MAC0422-2011/xv6$ ls /proc
1      1376  1602  1683  232   3907  690   buddyinfo    key-users    self
10     14     1606  1685  24     4     695   bus          kmsg        slabinfo
1006   15     1619  1687  249   4004  696   cgroups     kpagecount  softirqs
1016   1506   1624  1695  250   4005  7     cmdline     kpageflags stat
1030   1510   1625  17     27     4006  762   cpuinfo     latency_stats swaps
1039   1528   1627  1703  28     41     788   crypto      loadavg     sys
10856  1562   1629  1705  29     42     8     devices     locks       sysrq-trigger
1091   1565   1631  1707  3      43     852   diskstats  mdstat      sysvipc
10927  1566   1633  1716  30     44     856   dma         meminfo    timer_list
1097   1569   1644  18     309   45     862   driver      misc       timer_stats
11     1576   1645  19     31     46     863   execdomains modules   tty
1120   1578   1655  2      312   5     865   fb          mounts   uptime
12     1586   1658  20     32     518   866   filesystems mpt      version
1240   1589   1659  21     3460  6     870   fs          mtrr      version_signature
1278   1591   1660  22     3461  625   871   interrupts net      vmallocinfo
1282   1592   1671  22035 35     653   889   iomem      pagetypeinfo vmmemctl
1286   1594   1674  2240  3518  673   9     ioports    partitions vmstat
1292   1596   1676  23     37     679   935   irq       sched_debug zoneinfo
13     1598   1679  230   3771  686   acpi      kallsyms  schedstat
1340   16     1680  231   38     688   asound    kcore     scsi
```

Processo vs. “thread”

- ❑ Thread: separate streams of execution that share the same address space
- ❑ Process != Thread
 - One process can have multiple threads
 - Threads communicate more efficiently

“thread”: fio

Linux: tasks!

- Linux uses a neutral term: **tasks**
 - Tasks represent both processes and threads
 - When processes trap into the kernel, they share the Linux kernel's address space → kernel threads

Um processo possui pelo menos 1 thread!

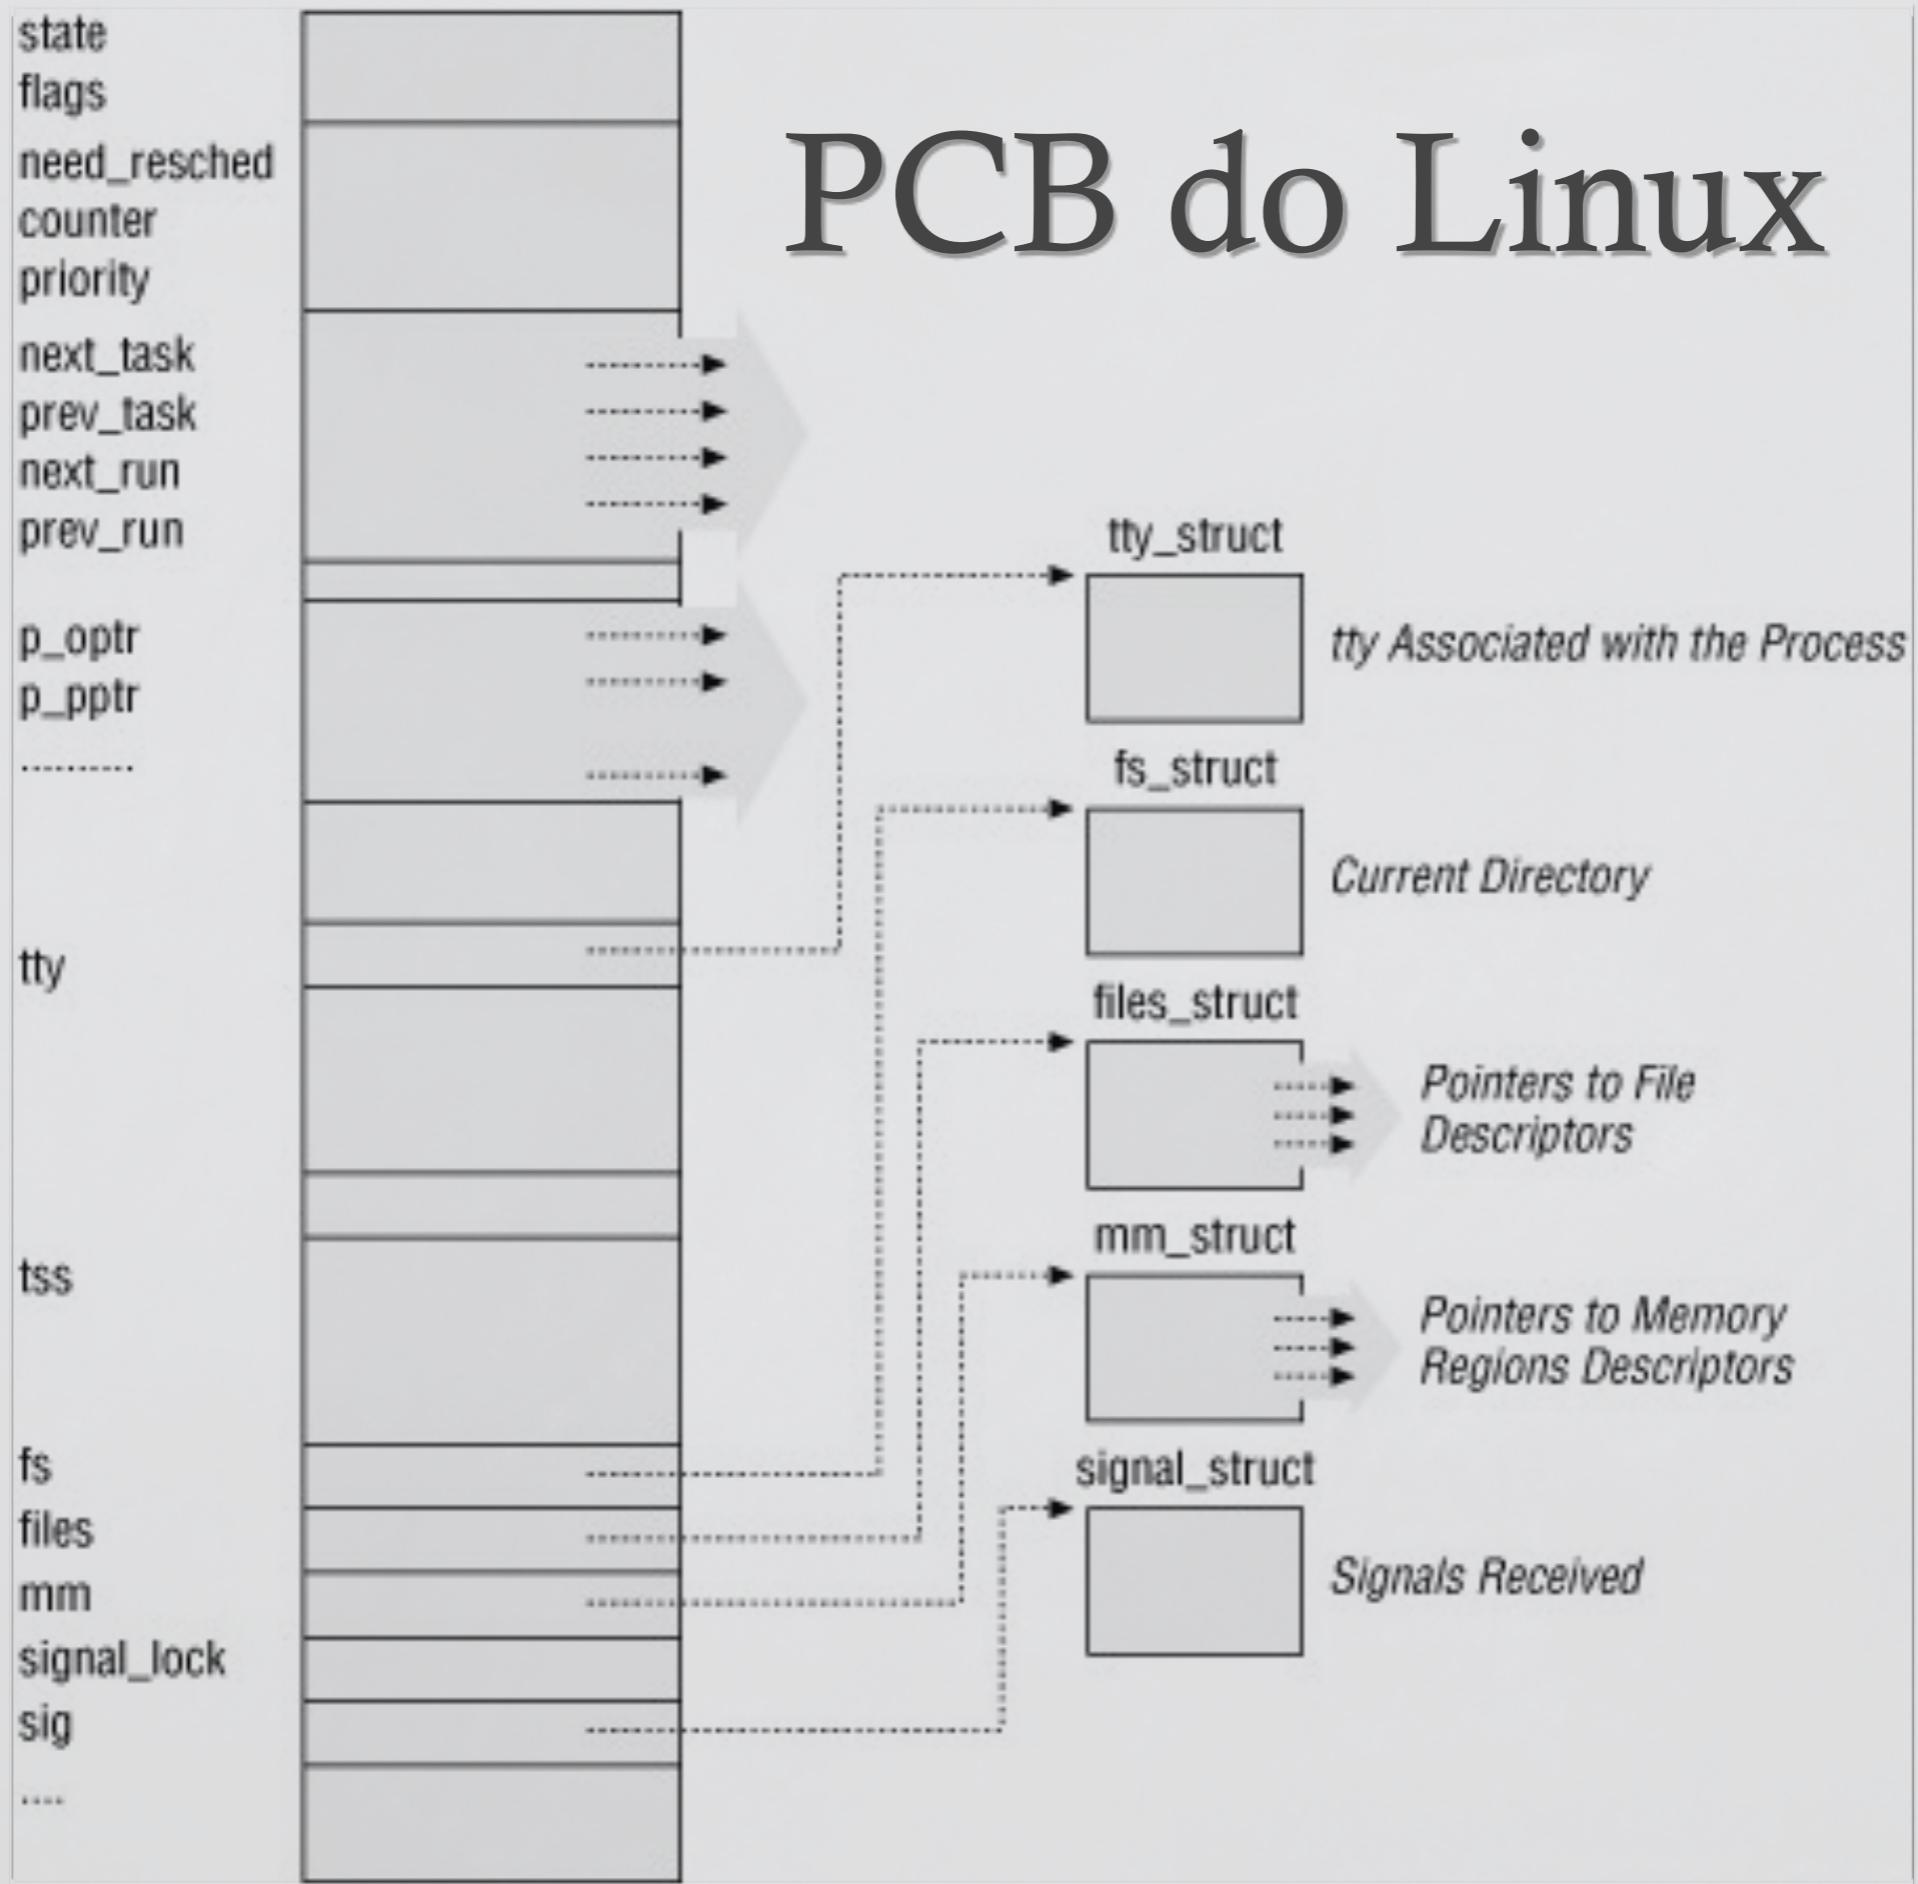
Como criar threads ?

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void* thread_fn(void *arg)
{
    int id = (int)arg;
    printf("thread %d runs\n", id);
    return NULL;
}

int main()
{
    pthread_t t1, t2;
    pthread_create(&t1, NULL, thread_fn, (void*)1);
    pthread_create(&t2, NULL, thread_fn, (void*)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    return 0;
}
```

PCB do Linux



Estados de um processo em Linux

- ❑ **state:** what state a process is in
 - **TASK_RUNNING** - the thread is running on the CPU or is waiting to run
 - **TASK_INTERRUPTIBLE** - the thread is sleeping and can be awoken by a signal (EINTR)
 - **TASK_UNINTERRUPTIBLE** - the thread is sleeping and cannot be awakened by a signal
 - **TASK_STOPPED** - the process has been stopped by a signal or by a debugger
 - **TASK_TRACED** - the process is being traced via the `ptrace` system call
- ❑ **exit_state:** how a process exited
 - **EXIT_ZOMBIE** - the process is exiting but has not yet been waited for by its parent
 - **EXIT_DEAD** - the process has exited and has been waited for

Identificadores de processos

- process ID: `pid`
- thread group ID: `tgid`
 - `pid` of first thread in process
 - `getpid()` returns this ID, so all threads in a process share the same process ID
- many system calls identify a process by its PID
 - Linux kernel uses `pidhash` to efficiently find processes by pids
 - see `include/linux/pid.h`, `kernel/pid.c`

getpid()

- Processos do usuário não podem executar operações privilegiadas por si próprios
 - Exemplo: getpid()
- Processos devem pedir ao SO para realizar tais chamadas através de “**system calls**”
- O SO é responsável por validar os parâmetros de cada um destes pedidos.

System calls (“syscalls”)

- Carregar registradores ou variáveis com parâmetros necessários
 - Instrução INT (“trap”)
 - Gera uma interrupção (de software)
 - Mudança automática para **modo “kernel”**
- Parâmetros dizem ao S.O. o que fazer
 - No seu término
 - “Retorna” como se fosse uma chamada à função
 - Força o retorno ao **modo “usuário”**

System calls: linux

Linux: *getpid()* in the standard C library looks like this:

```
0x400b65e0 <getpid+0>:    mov    $0x14,%eax  
0x400b65e5 <getpid+5>:    int    $0x80  
0x400b65e7 <getpid+7>:    cmp    $0xfffff001,%eax  
0x400b65ec <getpid+12>:   jae    0x400b65ef <getpid+15>  
0x400b65ee <getpid+14>:   ret  
0x400b65ef <getpid+15>:   ... ; Error handling  
0x400b660c <getpid+44>:   jmp    0x400b65ee <getpid+14>
```

0x80 (128) is the system call interrupt number, 0x14 (20) means *getpid*.

Return value of the system call is stored in the `%eax` register.

In Linux: an unsigned number larger than 0xfffff001 means error.

In 2's complement, it means a negative number in the range -4095 – -1.

The code above is called the C-library glue to system call.

So that the users don't need to know how to make INTs themselves.

Resumo

- A program calls the C-library function like `getpid()`.
- `getpid()` puts the service number (20) into the register `%eax`, and generates a software interrupt using `int $0x80`.
- The interrupt vector specifies that the kernel function `system_call` should be executed in kernel mode.
- `system_call` checks that the service number is valid, and looks up a table to find and execute the implementation function `sys_getpid()`.
- `sys_getpid()` finds the PID and puts the result to `%eax` and returns to `system_call`.
- `system_call` uses `IRET` to get back to the user program, and the CPU falls back to user mode.
- `getpid()` checks for error and returns to the user program.

Cadeia de chamadas para fork() em Linux

- ❑ libc fork()
- ❑ system_call (arch/i386/kernel/entry.S)
- ❑ sys_clone() (arch/i386/kernel/process.c)
- ❑ do_fork() (kernel/fork.c)
- ❑ copy_process() (kernel/fork.c)
 - ❑ p = dup_task_struct(current) // shallow copy
 - ❑ copy_* // copy point-to structures
 - ❑ copy_thread () // copy stack, regs, and eip
 - ❑ wake_up_new_task() // set child runnable

Cadeia de chamadas para exit() em Linux

- ❑ `libc exit(code)`
- ❑ `system_call (arch/i386/kernel/entry.S)`
- ❑ `sys_exit() (kernel/exit.c)`
- ❑ `do_exit() (kernel/exit.c)`
- ❑ `exit_*() // free data structures`
- ❑ `exit_notify() // tell other processes we exit`
 - `// reparent children to init`
 - `// EXIT_ZOMBIE`
 - `// EXIT_DEAD`

Exercício

- Em Linux, o número da interrupção que dá acesso à syscalls é a 0x80 (\$128).
 - E em xv6 ?
 - No Windows ?
- Usando o código-fonte do xv6 ou o debugger (gdb), liste todas as chamadas em cascata decorrente das syscalls fork() e exit().