



MAC422 Sistemas Operacionais

Prof. Marcel P. Jackowski

Aula #8

Threads e sincronização

fork()

```
// Create a new process copying p as the parent.  
// Sets up stack to return as if from system call.  
// Caller must set state of returned proc to RUNNABLE.  
int  
fork(void)  
{  
    int i, pid;  
    struct proc *np;  
    // Allocate process.  
    if((np = allocproc()) == 0)  
        return -1;  
  
    // Copy process state from p.  
    if((np->pgdir = copyuvm(proc->pgdir, proc->sz)) == 0){  
        kfree(np->kstack);  
        np->kstack = 0;  
        np->state = UNUSED;  
        return -1;  
    }  
    np->sz = proc->sz;  
    np->parent = proc;  
    *np->tf = *proc->tf;  
  
    // Clear %eax so that fork returns 0  
    // in the child.  
    np->tf->eax = 0;  
  
    for(i = 0; i < NOFILE; i++)  
        if(proc->ofile[i])  
            np->ofile[i] = filedup(proc->ofile[i]);  
    np->cwd = idup(proc->cwd);  
  
    pid = np->pid;  
    np->state = RUNNABLE;  
    safestrcpy(np->name, proc->name, sizeof  
(proc->name));  
    return pid;  
}
```

exit()

```
// Exit the current process. Does not return.  
// An exited process remains in the zombie state  
// until its parent calls wait() to find out it exited.  
void  
exit(void)  
{  
    struct proc *p;  
    int fd;  
  
    if(proc == initproc)  
        panic("init exiting");  
  
    // Close all open files.  
    for(fd = 0; fd < NOFILE; fd++){  
        if(proc->ofile[fd]){  
            fileclose(proc->ofile[fd]);  
            proc->ofile[fd] = 0;  
        }  
    }  
  
    iput(proc->cwd);  
    proc->cwd = 0;  
  
    acquire(&ptable.lock);  
  
    // Parent might be sleeping in wait().  
    wakeup1(proc->parent);  
  
    // Pass abandoned children to init.  
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){  
        if(p->parent == proc){  
            p->parent = initproc;  
            if(p->state == ZOMBIE)  
                wakeup1(initproc);  
        }  
    }  
  
    // Jump into the scheduler, never to return.  
    proc->state = ZOMBIE;  
    sched();  
    panic("zombie exit");  
}
```

wait()

```
// Wait for a child process to exit and return its pid.  
// Return -1 if this process has no children.  
int  
wait(void)  
{  
    struct proc *p; int havekids, pid;  
    acquire(&ptable.lock);  
    for(;;){  
        // Scan through table looking for zombie children.  
        havekids = 0;  
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {  
            if(p->parent != proc)  
                continue;  
            havekids = 1;  
            if(p->state == ZOMBIE) { // Found one.  
                pid = p->pid;  
                kfree(p->kstack); p->kstack = 0; freevm(p->pgdir);  
                p->state = UNUSED; p->pid = 0; p->parent = 0; p->name[0] = 0; p->killed = 0;  
                release(&ptable.lock);  
                return pid;  
            }  
        }  
  
        // No point waiting if we don't have any children.  
        if(!havekids || proc->killed){  
            release(&ptable.lock); return -1;  
        }  
  
        // Wait for children to exit. (See wakeup1 call in proc_exit.)  
        sleep(proc, &ptable.lock); //DOC: wait-sleep  
    }  
}
```

```

#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#define MAX 10
void print(int *v)
{
    int i;

    for(i=0; i<MAX; i++)
        printf("%d ", v[i]);
    printf("\n");
}

int main(void)
{
    int i, x;
    int *y = (int *) malloc (MAX * sizeof(int));
    for(i=0;i<10;i++) y[i] = i;
    pid_t pid = fork();
    if (pid == 0) {
        printf("I am in the child process, y=%x\n", y);
        sleep(2);
        /* print vector */
        print(y);
        free(y);
    } else {
        /* If fork() returns a positive number,
           we are in the parent process and pid is the
           child process id.
        */
        printf("I am in the parent process, y=%x\n", y);
        y[9] = 0;
        /* print vector */
        print(y);
        free(y);
        if (fork() > 0)
            fprintf(stderr, "Error, can't fork!\n");
        exit(1);
    }
    exit(0);
}

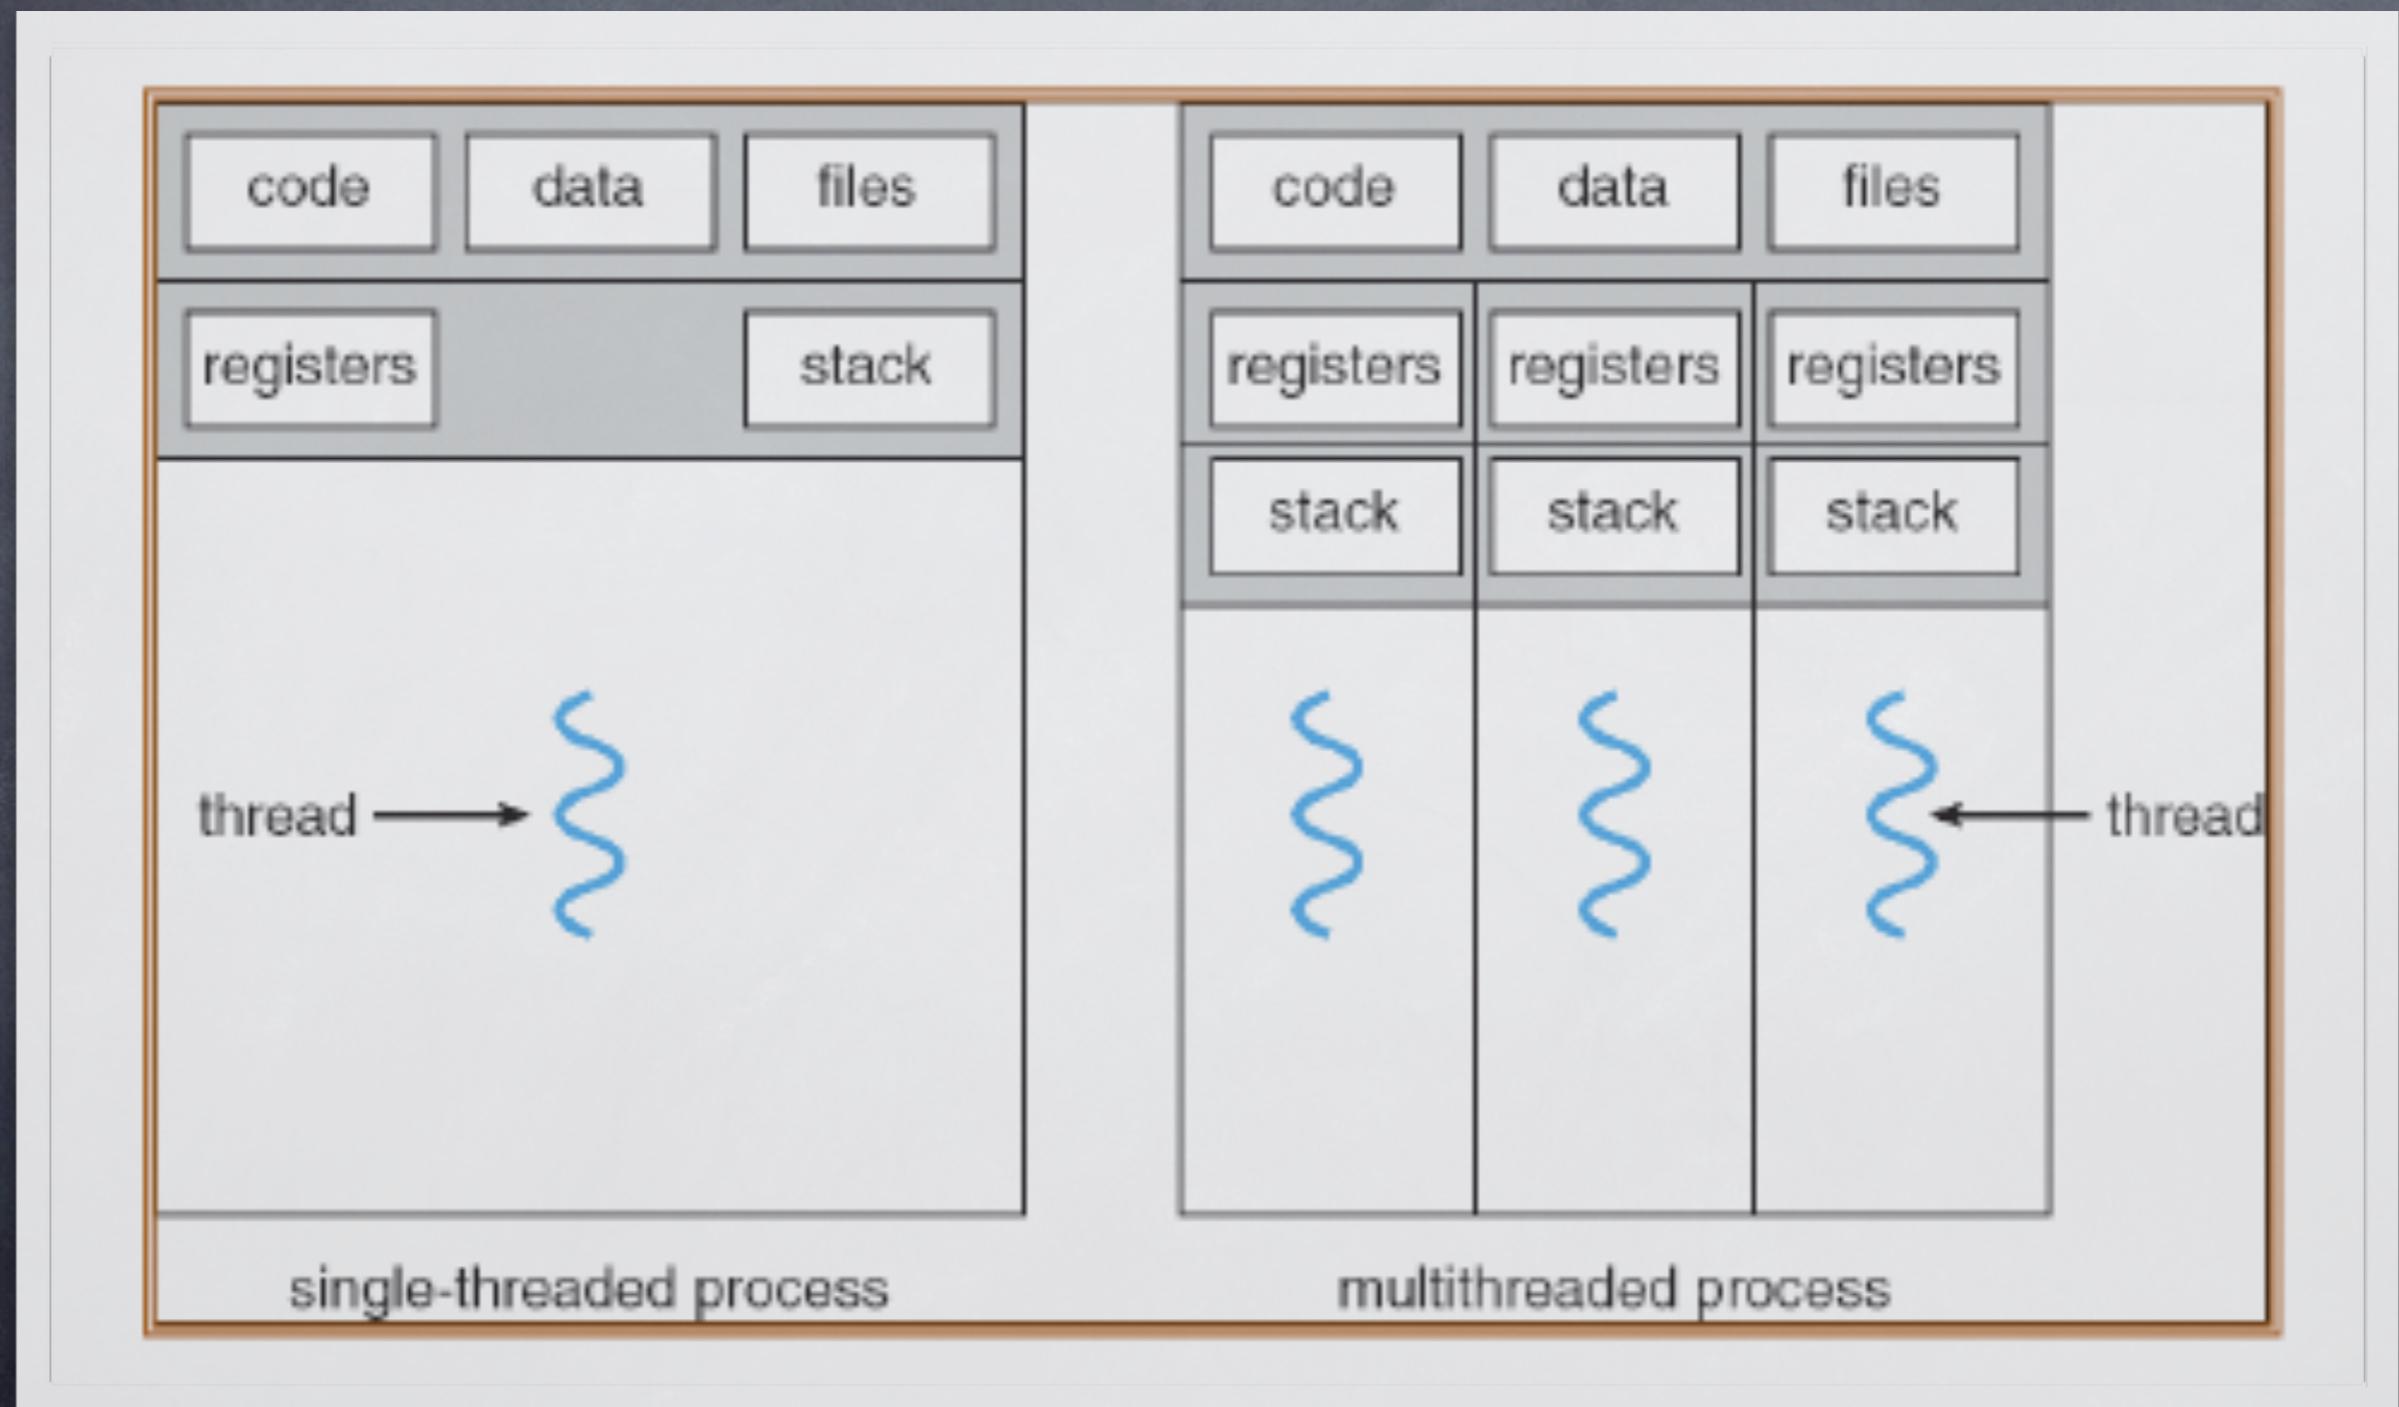
```

O que acontece com o espaço
alocado por malloc() ?

Threads

- **Threads:** separate streams of executions that share an address space
 - Allows one process to have multiple point of executions, can potentially use multiple CPUs
- **Thread control block (TCB)**
 - Program counter (EIP on x86)
 - Other registers
 - Stack

Single-threaded e multithreaded



Por quê threads ?

- Express **concurrency**

- Web server (multiple requests), Browser (GUI + network I/O + rendering), ...

```
for(;;) {  
    struct request *req = get_request();  
    create_thread(process_request, req);  
}
```

- Efficient communication

- Using a separate process for each task can be heavyweight

Threads vs. Processes

- A thread has no data segment or heap
- A thread cannot live on its own, it must live within a process
- There can be more than one thread in a process, the first thread calls `main()` & has the process's stack
- Inexpensive creation
- Inexpensive context switching
- Efficient communication
- If a thread dies, its stack is reclaimed
- A process has code/data/heap & other segments
- A process has at least one thread
- Threads within a process share code/data/heap, share I/O, but each has its own stack & registers
- Expensive creation
- Expensive context switching
- Interprocess communication can be expressive
- If a process dies, its resources are reclaimed & all threads die

Utilizando threads

- Through thread library
 - E.g. pthread, Win32 thread
- Common operations
 - create/terminate
 - suspend/resume
 - priorities and scheduling
 - synchronization

Pthreads

- `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void *(*start_routine)(void*), void *arg);`
 - Create a new thread to run `start_routine` on `arg`
 - `thread` holds the new thread's id
 - Can be customized via `attr`
- `int pthread_join(pthread_t thread, void **value_ptr);`
 - Wait for `thread` termination, and retrieve return value in `value_ptr`
- `void pthread_exit(void *value_ptr);`
 - Terminates the calling thread, and returns `value_ptr` to threads waiting in `pthread_join`

Exemplo de criação de threads

```
void* thread_fn(void *arg)
{
    int id = (int)arg;
    printf("thread %d runs\n", id);
    return NULL;
}

int main()
{
    pthread_t t1, t2;
    pthread_create(&t1, NULL, thread_fn, (void*)1);
    pthread_create(&t2, NULL, thread_fn, (void*)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    return 0;
}
```

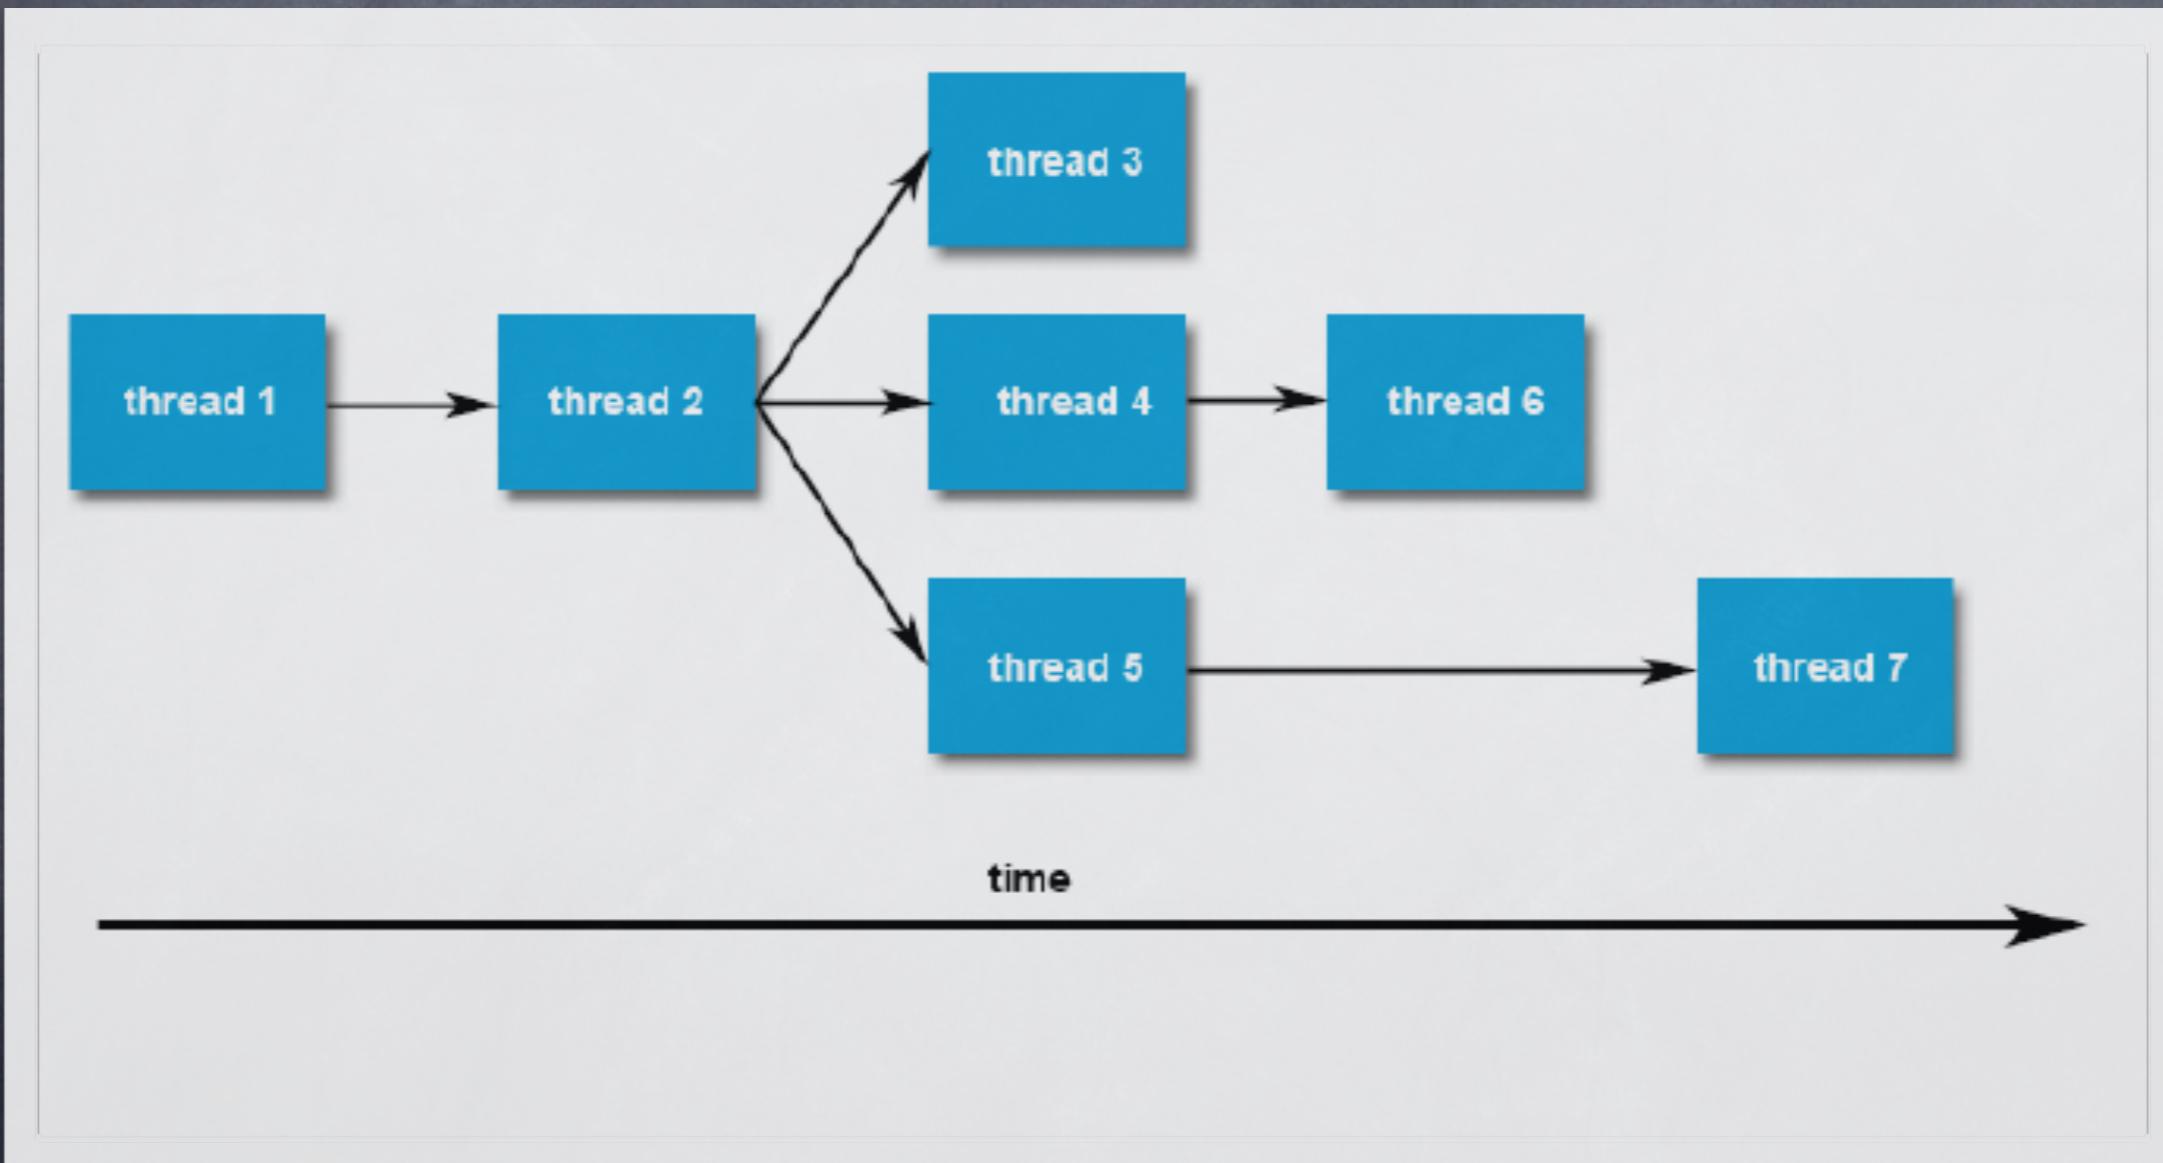
```
$ gcc -o threads threads.c -Wall -lpthread
$ threads
thread 1 runs
thread 2 runs
```

One way to view threads: function calls, except caller doesn't wait for callee; instead, both run concurrently

Exemplo de pthreads

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 5
void *PrintHello(void *threadid)
{
    int tid;
    tid = (int)threadid;
    printf("Hello World! It's me, thread #%-d!\n", tid);
    pthread_exit(NULL);
}
int main(int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc, t;
    for(t=0;t<NUM_THREADS;t++){
        printf("In main: creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc){
            printf("ERROR: return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

Threads podem criar outros threads

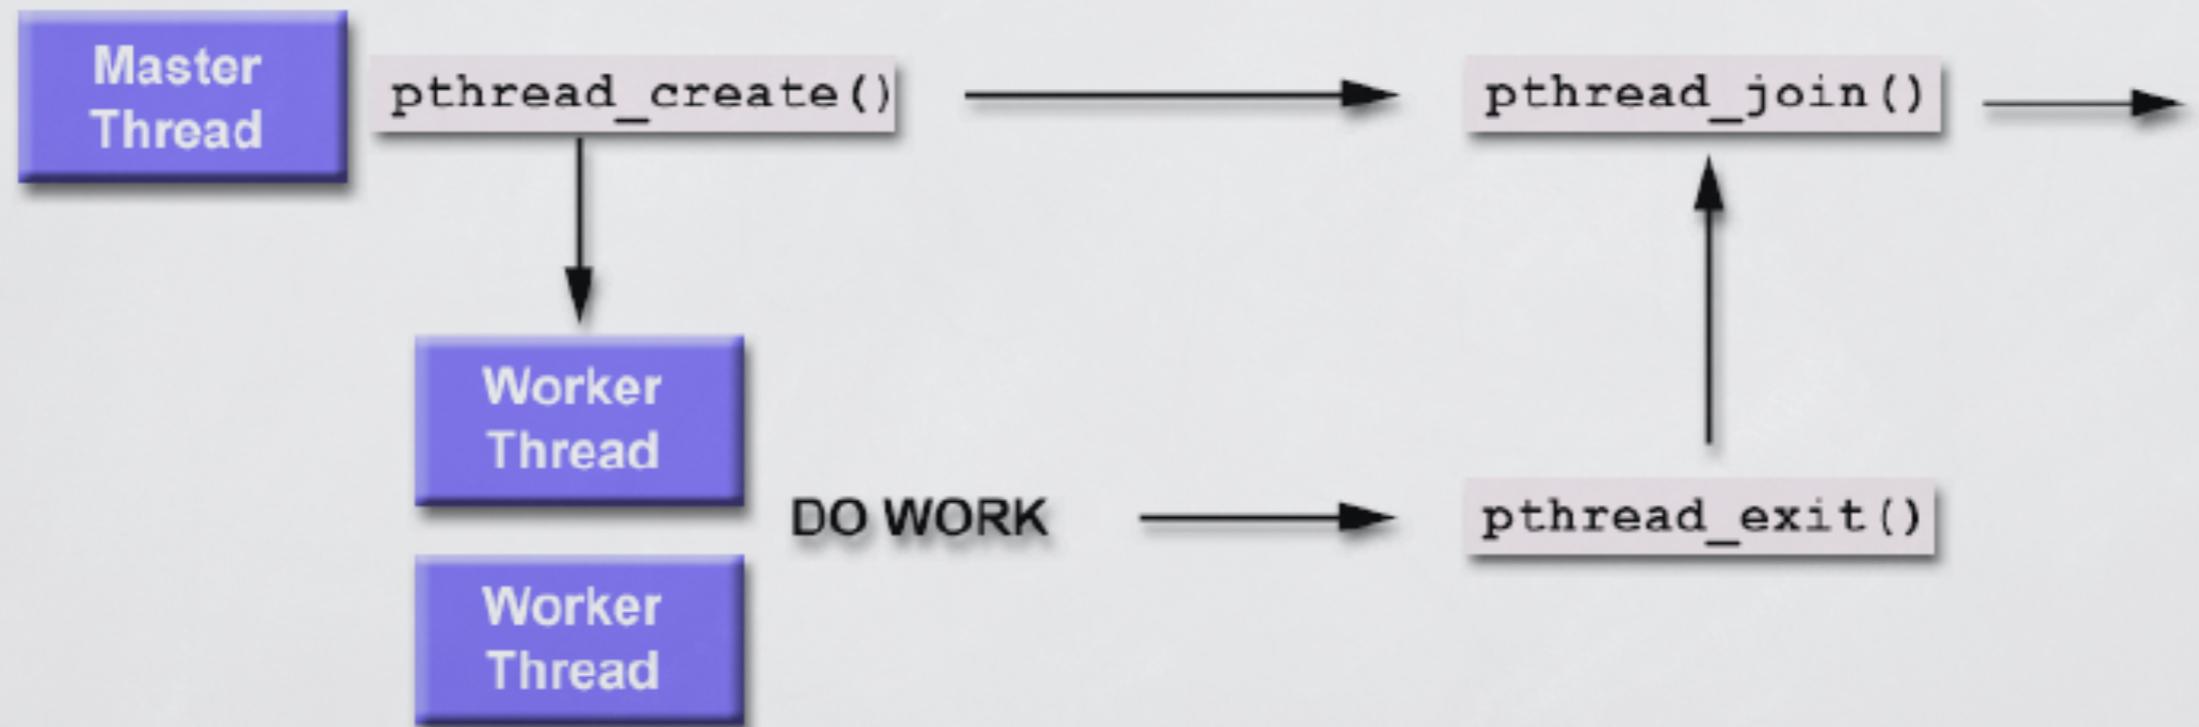


Quem é escalonado primeiro ?

Terminando threads

- No término de sua rotina de execução
- A thread chama `pthread_exit()`
- A thread é cancelada por outra com o comando `pthread_cancel()`
- O processo inteiro é cancelado através de uma chamada a `exit()` ou `exec()`.

Sincronizando o trabalho



```
int main(int argc, char *argv[])
{
    pthread_t thread[NUM_THREADS];
    pthread_attr_t attr;
    int rc, t;
    void *status;

    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
    for(t=0;t<NUM_THREADS;t++)
    {
        printf("Creating thread %d\n", t);
        rc = pthread_create(&thread[t], &attr, BusyWork, NULL);
        if (rc)
            printf("ERROR; return code from pthread_create() is %d\n", rc);
        exit(-1);
    }
    pthread_attr_destroy(&attr);
    for(t=0;t<NUM_THREADS;t++)
    {
        rc = pthread_join(thread[t], &status);
        if (rc)
            printf("ERROR return code from pthread_join() is %d\n", rc);
        exit(-1);
        printf("Completed join with thread %d status= %ld\n",t,(long)status);
    }
    pthread_exit(NULL);
}
```

Continuação

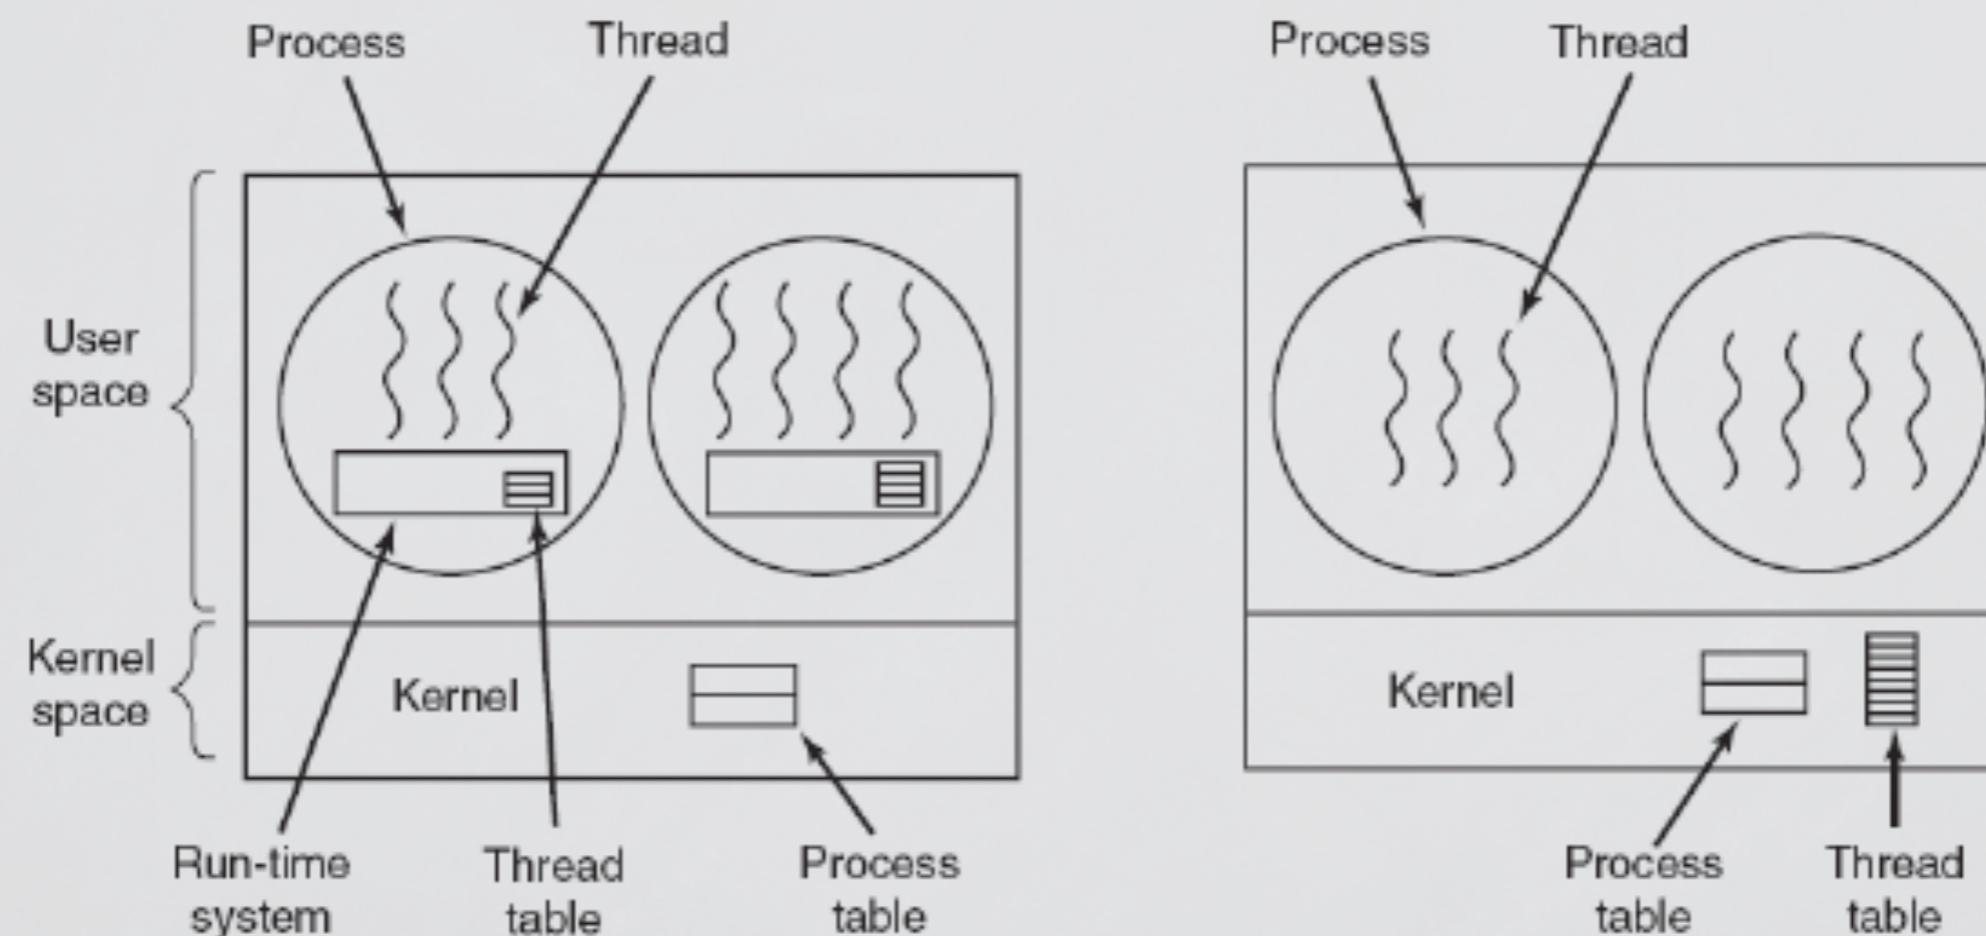
```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 3

void *BusyWork(void *null)
{
    int i;
    double result=0.0;
    for (i=0; i<1000000; i++)
    {
        result = result + (double) random();
    }
    printf("Thread result = %e\n", result);
    pthread_exit((void *) 0);
}
```

Modelos de multithreading

- Where to support threads?
- User threads: thread management done by user-level threads library; kernel knows nothing
- Kernel threads: threads directly supported by the kernel
 - Virtually all modern OS support kernel threads

User vs. kernel threads



Example from Tanenbaum, Modern Operating Systems 3 e,
(c) 2008 Prentice-Hall, Inc. All rights reserved. 0-13-6006639

User threads vs. kernel threads

- Pros: fast, no system call for creation, context switch
- Cons: kernel doesn't know → one thread blocks, all threads in the process blocks
- Cons: slow, kernel does creation, scheduling, etc
- Pros: kernel knows → one thread blocks, schedule another

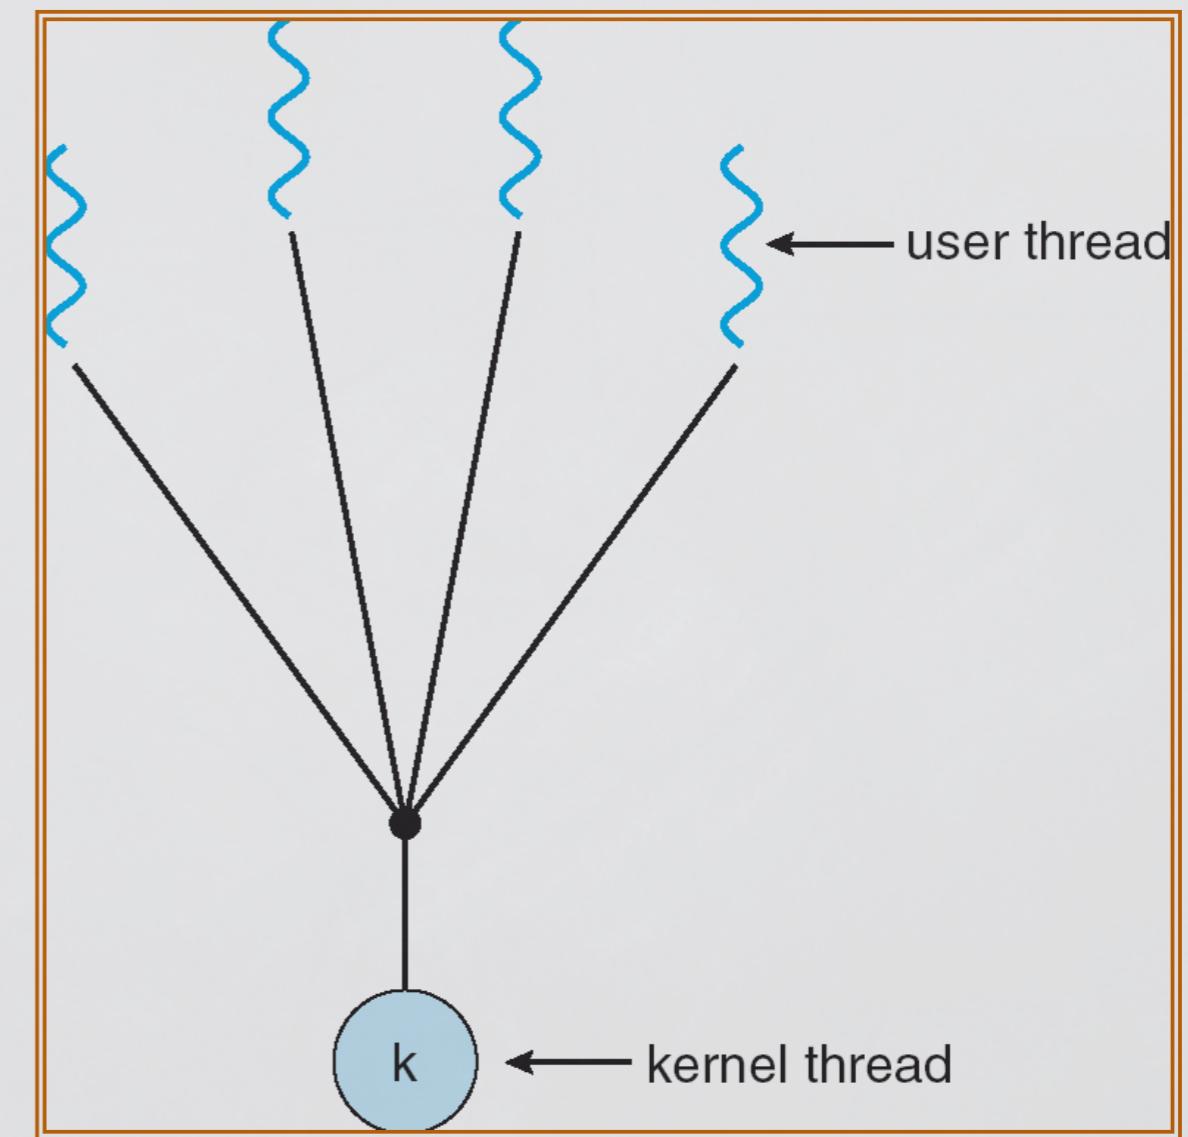
No free lunch!

Implementação de threads

- A thread library must map user threads to kernel threads
- Big picture:
 - kernel thread: physical concurrency, how many cores?
 - User thread: application concurrency, how many tasks?
- Different mappings exist, representing different tradeoffs
 - Many-to-One: many user threads map to one kernel thread, i.e. kernel sees a single process
 - One-to-One: one user thread maps to one kernel thread
 - Many-to-Many: many user threads map to many kernel threads

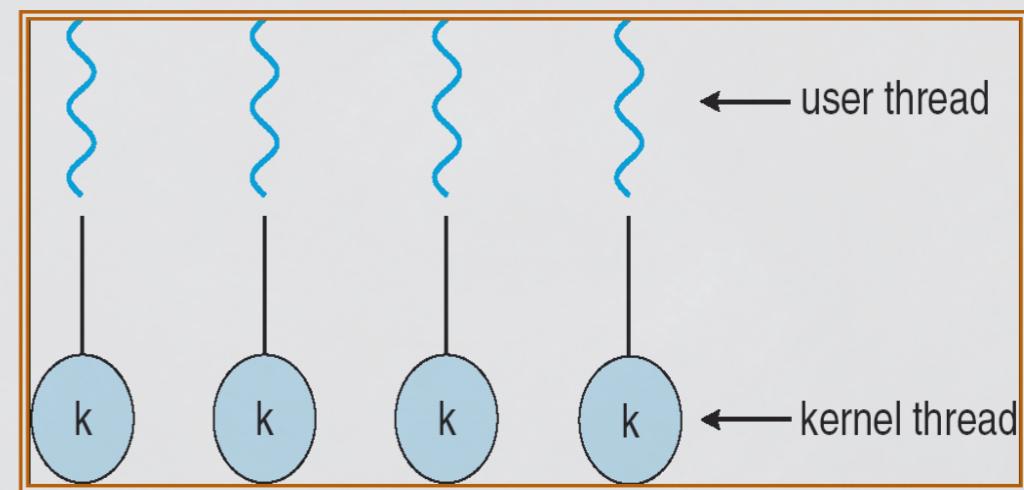
“Many-to-one”

- Many user-level threads map to one kernel thread
- Pros
 - **Fast**: no system calls required
 - **Portable**: few system dependencies
- Cons
 - **No parallel execution of threads**
 - All thread block when one waits for I/O



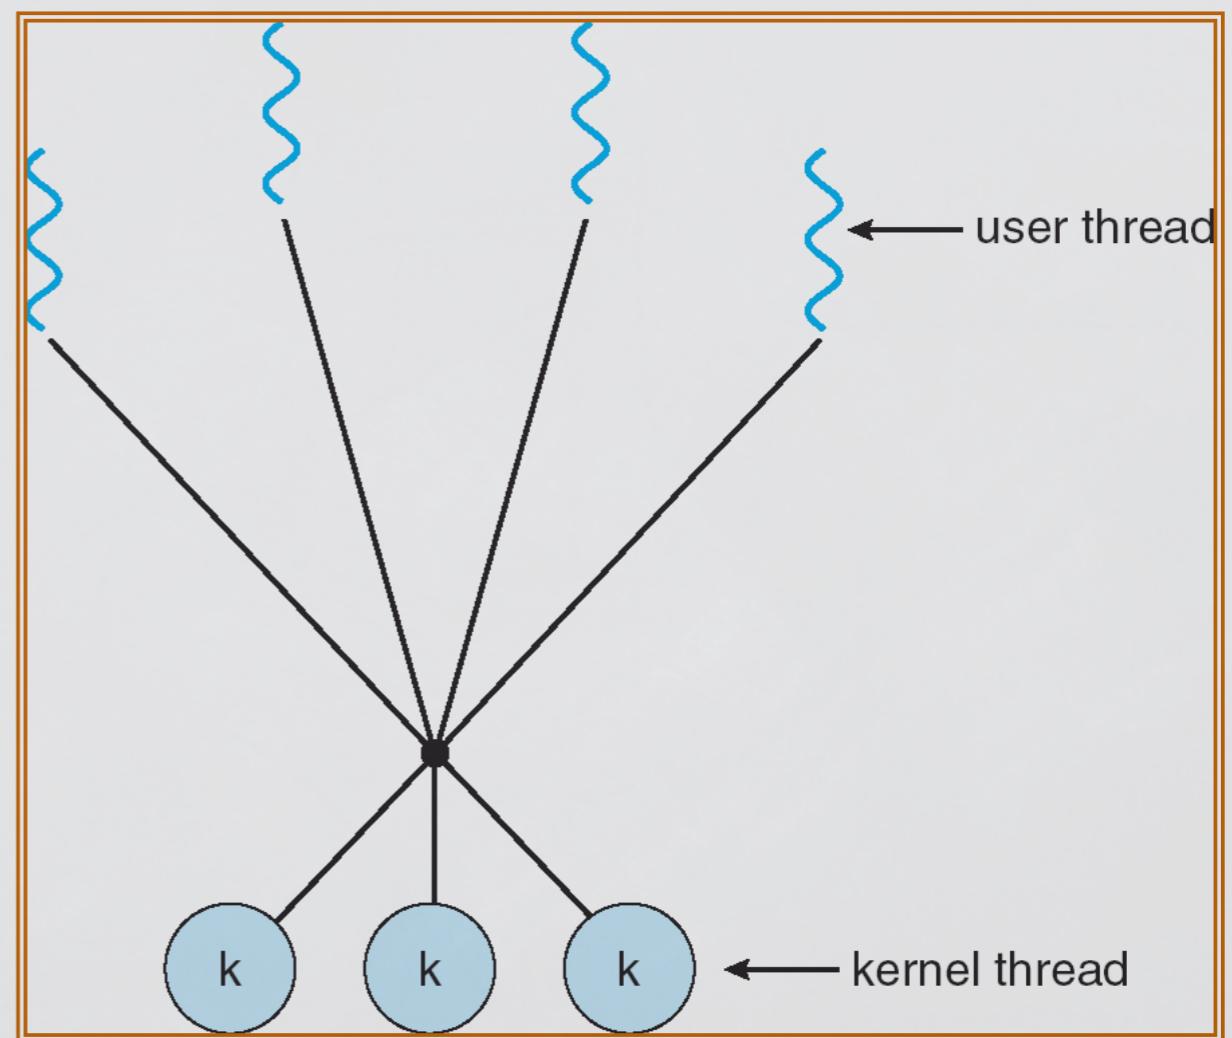
“One-to-one”

- One user-level thread maps to one kernel thread
- Pros: **more concurrency**
 - When one blocks, others can run
 - Better multicore or multiprocessor performance
- Cons: **expensive**
 - Thread operations involve kernel
 - Thread need kernel resources



“Many-to-many”

- Many user-level threads map to many kernel threads ($U \geq K$)
 - Supported some versions of BSD, and Windows
- Pros: **flexible**
 - OS creates kernel threads for physical concurrency
 - Applications creates user threads for application concurrency
- Cons: **complex**
 - Most programs use 1:1 mapping anyway



Threads em Linux

- Unidade de escalonamento: “tasks”
- Criação através do *syscall* “`clone()`”
- Utilização do modelo “kernel threads”
 - One-to-one
- Exemplos de informações que são compartilhadas entre threads:
 - Espaço em memória
 - Descritores de arquivos
 - Gerenciadores de sinais

Implementação em Windows™

- Processos não são escalonáveis.
- Cada processo possui inicialmente 1 thread.
 - Threads são escalonáveis
- Novos threads são criados através da *syscall*
 - `CreateThread()`
 - Passa-se o endereço da função ou código a ser executado.
- Fibras (“*fibers*”)

Threads em Java

- Podemos usar herança e criamos uma classe que estende a classe **Thread**
- Devemos redefinir o método **run()**:

```
class Tarefa1 extends Thread {  
    public void run() {  
        for(int i=0; i<1000; i++)  
            System.out.println("Usando herança");  
    }  
}
```

```
Thread threadComHeranca = new Tarefa1();  
threadComHeranca.start();
```

Threads em Java

- A interface **Runnable** nos obriga a implementar o método **run()**:

```
class Tarefa2 implements Runnable {  
    public void run() {  
        for(int i=0; i<1000; i++)  
            System.out.println("Usando Runnable");  
    }  
}
```

```
Thread threadComRunnable = new Thread(new Tarefa2());  
threadComRunnable.start();
```

“Thread pool”

- Problem:
 - Creating a thread for each request: **costly**
 - And, the created thread exits after serving a request
 - More user request → More threads, **server overload**
- Solution: **thread pool**
 - Pre-create a number of threads waiting for work
 - Wake up thread to serve user request --- **faster than thread creation**
 - When request done, don't exit --- go back to pool
 - **Limits the max number of threads**

Transação bancária

```
int balance = 1000;

int main()
{
    pthread_t t1, t2;
    pthread_create(&t1, NULL, deposit, (void*)1);
    pthread_create(&t2, NULL, withdraw, (void*)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("all done: balance = %d\n", balance);
    return 0;
}
```

```
void* deposit(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i)
        ++balance;
}

void* withdraw(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i)
        --balance;
}
```

Possíveis resultados

```
$ gcc -Wall -lpthread -o bank bank.c
```

```
$ bank
```

```
all done: balance = 0
```

```
$ bank
```

```
all done: balance = 140020
```

```
$ bank
```

```
all done: balance = -94304
```

```
$ bank
```

```
all done: balance = -191009
```

Por quê ?

Olhando mais a fundo...

```
$ objdump -d bank
```

```
...
```

```
08048464 <deposit>:
```

```
... // ++ balance
```

```
8048473: a1 80 97 04 08    mov 0x8049780,%eax
```

```
8048478: 83 c0 01          add $0x1,%eax
```

```
804847b: a3 80 97 04 08    mov %eax,0x8049780
```

```
...
```

```
0804849b <withdraw>:
```

```
... // -- balance
```

```
80484aa: a1 80 97 04 08    mov 0x8049780,%eax
```

```
80484af: 83 e8 01          sub $0x1,%eax
```

```
80484b2: a3 80 97 04 08    mov %eax,0x8049780
```

Possível escalonamento

	CPU 0	CPU 1
		balance: 1000
	mov 0x8049780,%eax	eax0: 1000
	add \$0x1,%eax	eax0: 1001
	mov %eax,0x8049780	balance: 1001
time		eax1: 1001
		sub \$0x1,%eax
		mov %eax,0x8049780
		balance: 1000
	One deposit and one withdraw, balance unchanged. Correct	

Outro escalonamento

	CPU 0	CPU 1
		balance: 1000
	mov 0x8049780,%eax	eax0: 1000
	add \$0x1,%eax	eax0: 1001
		eax1: 1000
	mov %eax,0x8049780	balance: 1001
time		eax1: 999
		balance: 999
		mov 0x8049780,%eax
		sub \$0x1,%eax
		mov %eax,0x8049780

One deposit and one withdraw,
balance becomes less. Wrong!

“Race condition”

- Definition: a timing dependent error involving shared state
- Can be very bad
 - “non-deterministic:” don’t know what the output will be, and it is likely to be different across runs
 - Hard to detect: too many possible schedules
 - Hard to debug: “heisenbug,” debugging changes timing so hides bugs (vs “bohr bug”)

Como evitar as “race conditions”?

- ❑ **Atomic operations**: no other instructions can be interleaved, executed “as a unit” “all or none”, guaranteed by hardware
- ❑ A possible solution: create a super instruction that does what we want atomically
 - add \$0x1, 0x8049780
- ❑ Problem
 - Can't anticipate **every possible** way we want atomicity
 - Increases hardware complexity, **slows down** other instructions

```
// ++ balance  
mov 0x8049780,%eax  
add $0x1,%eax  
mov %eax,0x8049780
```

```
...  
  
// -- balance  
mov 0x8049780,%eax  
sub $0x1,%eax  
mov %eax,0x8049780
```

...

Sincronização por camadas

- ❑ Hardware provides simple **low-level atomic operations**, upon which we can build **high-level, synchronization primitives**, upon which we can implement critical sections and build correct multi-threaded/multi-process programs

Properly synchronized application

High-level synchronization
primitives

Hardware-provided low-level
atomic operations

Primitivas de sincronização

- Low-level atomic operations
 - On uniprocessor, disable/enable interrupt
 - On x86, aligned load and store of words
 - Special instructions:
 - test-and-set (TSL), compare-and-swap (XCHG)
- High-level synchronization primitives
 - Lock
 - Semaphore
 - Monitor

Tarefa de Casa

- Leitura sobre o xv6:
 - Veja como exec() funciona no código fonte
 - sysfile.c
 - exec.c
- Documentação sobre o xv6:
 - exec.pdf
 - lock.pdf